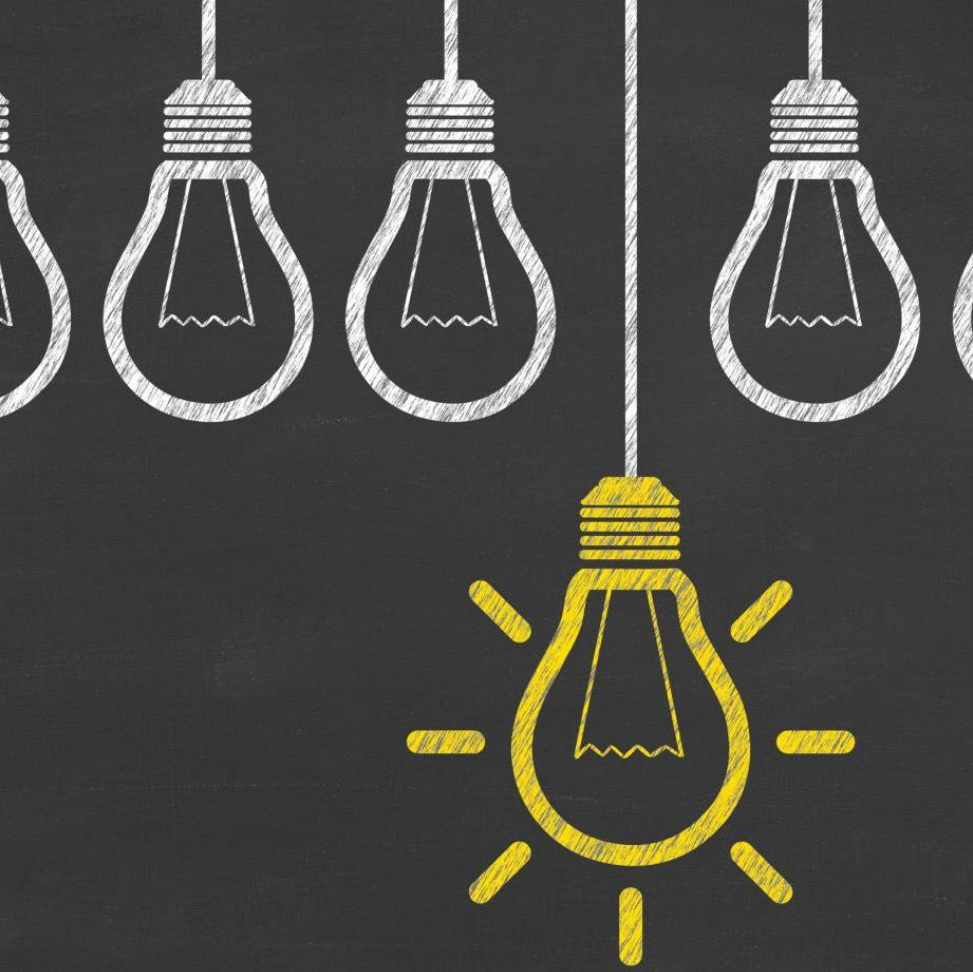




# Language Models

Liad Magen

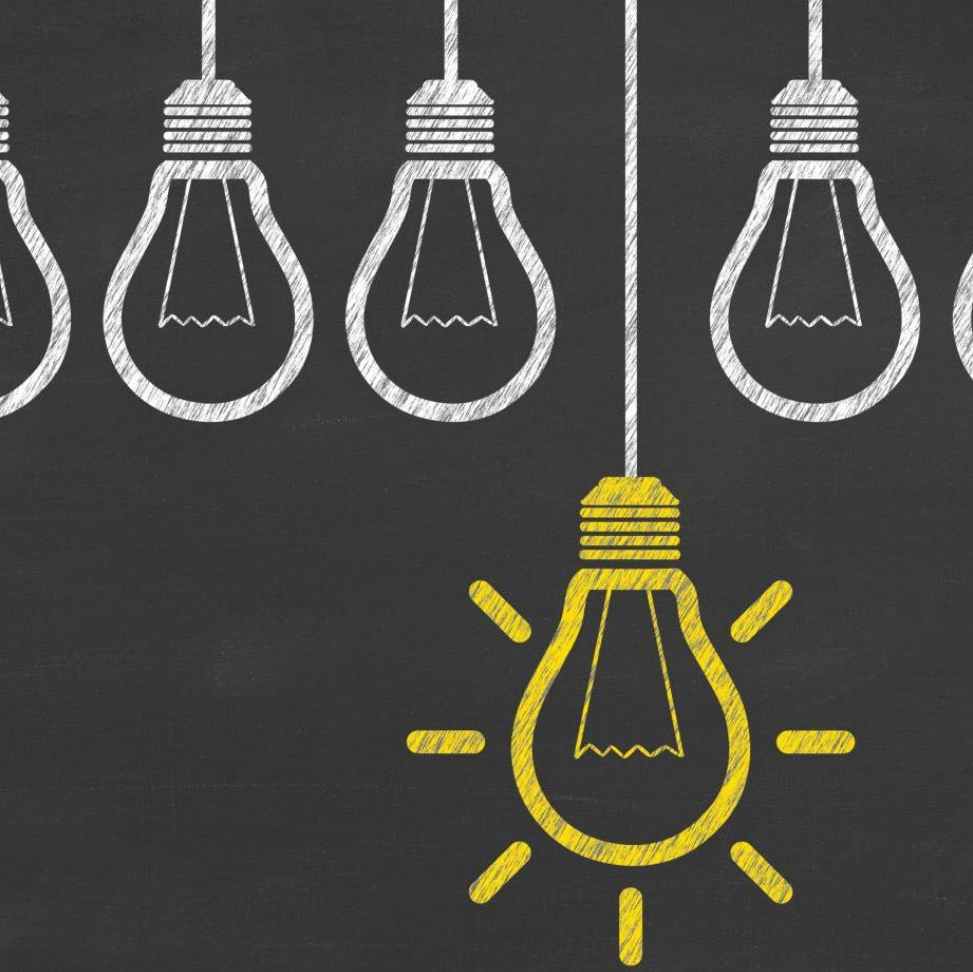


# Sequences

---

Let's play a game:  
what is the next word?

I'm gonna make him an offer he can't ...

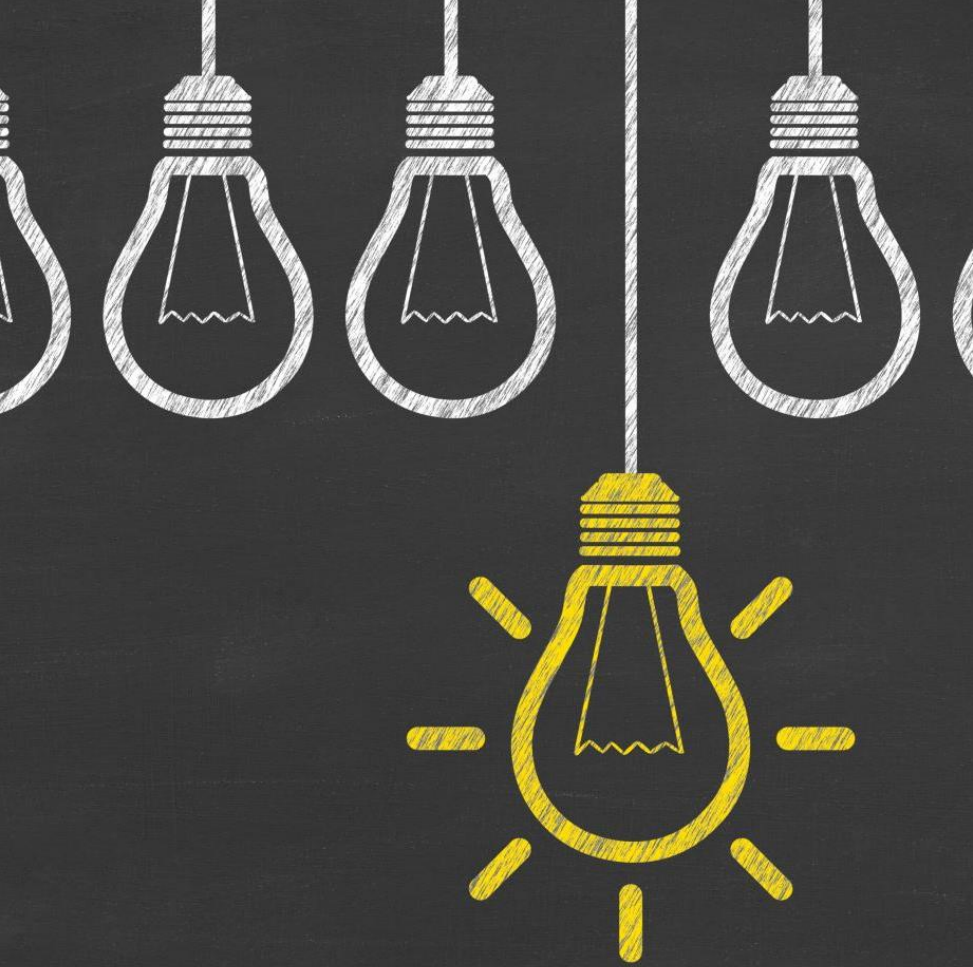


# Sequences

---

Let's play a game:  
what is the next word?

May the force be with ...



# Sequences

---

Let's play a game:  
what is the next word?

There's no place like ...

# Language Modeling

Giving a sequence of words ( $x_1 \dots x_{i-1}$ ),  
compute the probability distribution of the  
next word.

$$p(x_i | x_1, \dots, x_{i-1})$$

P( Home | There's no place like )

# Language Model

Can also assign a *probability score* for a sequence:

$$\begin{aligned} p(x_1, \dots, x_n) = & p_{LM}(x_1 | *S*, *S*) \\ & \times p_{LM}(x_2 | *S*, x_1) \\ & \times p_{LM}(x_3 | x_1, x_2) \\ & \times p_{LM}(x_4 | x_2, x_3) \\ & \dots \\ & \times p_{LM}(x_n | x_{n-2}, x_{n-1}) \end{aligned}$$

P( Their are two examples )

P( There are two examples )

State Warriors 96 to 82 in the 4th quarter.

NBA · Today

Q4 - 10:39



Warriors

96



Spurs

82

NBA

Watch on: ESPN



Search results



Highlights

NBA standin

When is



your

memorial day

easter

1 2 3 4 5 6 7 8 9 0  
q w e r t y u i o p

a s d f g h j k l

↑ z x c v b n m ✕

?123 , 😊 . ➡

# Language Model

- Very useful: also used in *Speech Recognition*, *Machine Translation*, etc.
- **Note:** Doesn't have to be over natural language. Ideas for other usage examples?

How is it  
calculated?

- Markov Assumption:  
 $X_i$  depends only on the preceding  $n-1$  words
- n-gram Language Models:

$$p(x_i | x_1, \dots, x_{i-1}) = p(x_i / x_{i-n+1}, \dots, x_{i-1})$$



# n-gram Language Models: Example

Suppose  $n=4$ :

~~When we collaborate with each other, we~~ can achieve great \_\_\_\_\_

$$P(\mathbf{w} \mid \text{can achieve great}) = \frac{\text{count}(\text{can achieve great } \mathbf{w})}{\text{count}(\text{can achieve great})}$$

# n-gram Language Models: Sparsity issues

Suppose  $n=4$ :

~~When we collaborate with each other, we can achieve great~~ \_\_\_\_\_

$$P(\mathbf{w} \mid \text{can achieve great}) = \frac{\text{count}(\text{can achieve great } \mathbf{w})}{\text{count}(\text{can achieve great})}$$

What if it never appears  
and  $P(W) = 0$  ?

What if this n-grams never  
occurs and the dominator  
is 0?

**Note:** The bigger  $n$  is, the worse our sparsity.  
Normally  $n$  won't be bigger than 5.

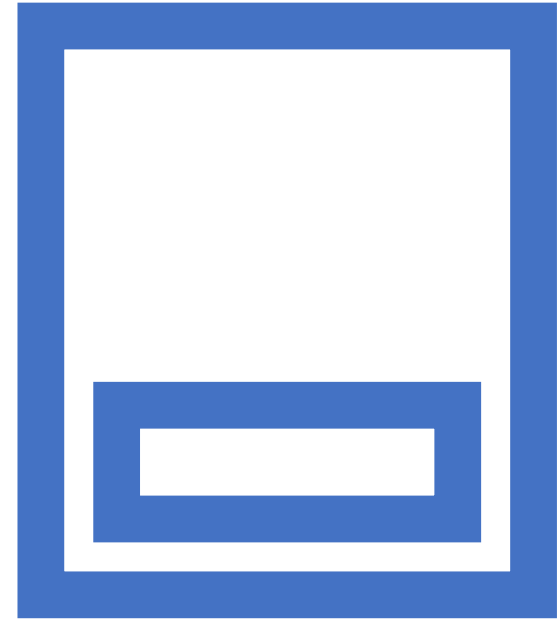
# Generating text with a n-gram Language Model

*Profit after financial \_\_\_\_\_*

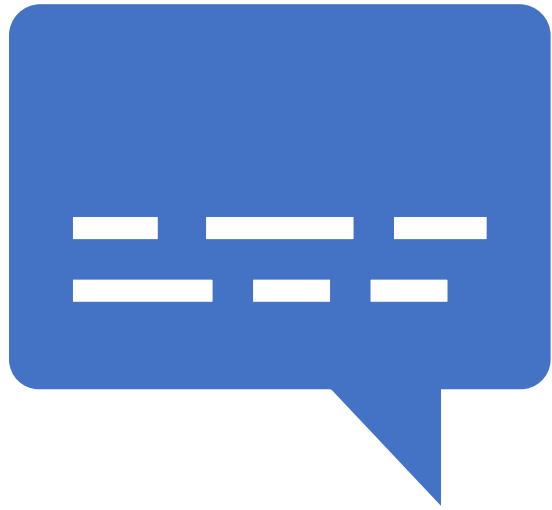
Sample from the probability distribution:

word	probability
income	0.035
crisis	0.022
support	0.031
report	0.032
with	0.0000001
run	0.0000001
...	...

Try it yourself



# Language Model



## Language Modeling:

Input: sequence of words:  $x_1, x_2, x_3 \dots x_n$

Output: *probability distribution* of the next word:  $P(x_{n+1} \mid x_n, x_{n-1} \dots x_1)$

# *Neural* Language Model

***Neural* Language Modeling** runs on a fixed-window, like n-gram:

Input: fixed-window sequence of last  $k$  words:

$$P(x_n \mid x_{n-1}, x_{n-2} \dots x_{n-k}) = \textit{softmax}(\textit{MLP}(x))$$

$$X = \textit{encode}(x_{n-1}, x_{n-2} \dots x_{n-k})$$

How do we encode the text?



# One-hot encoding

We have  $k$  elements in a vocabulary of size  $|V|$

Let's assume  $k=4$ ,  $|V| = 10$  and we want to encode( $x_1, x_2, x_3, x_4$ )

$V=\{A,B,C,D,E,F,G,H,I,J\}$

---

# One-hot-encoding of $k$ elements

A=[1,0,0,0,0,0,0,0,0,0]

B=[0,1,0,0,0,0,0,0,0,0]

C=[0,0,1,0,0,0,0,0,0,0]

D=[0,0,0,1,0,0,0,0,0,0]

E=[0,0,0,0,1,0,0,0,0,0]

F=[0,0,0,0,0,1,0,0,0,0]

G=[0,0,0,0,0,0,1,0,0,0]

H=[0,0,0,0,0,0,0,1,0,0]

I=[0,0,0,0,0,0,0,0,1,0]

J=[0,0,0,0,0,0,0,0,0,1]

**How should we encode (D, A, G, C) ?**



# One-hot-encoding of $k$ elements

A=[1,0,0,0,0,0,0,0,0,0]

B=[0,1,0,0,0,0,0,0,0,0]

C=[0,0,1,0,0,0,0,0,0,0]

D=[0,0,0,1,0,0,0,0,0,0]

E=[0,0,0,0,1,0,0,0,0,0]

F=[0,0,0,0,0,1,0,0,0,0]

G=[0,0,0,0,0,0,1,0,0,0]

H=[0,0,0,0,0,0,0,1,0,0]

I=[0,0,0,0,0,0,0,0,1,0]

J=[0,0,0,0,0,0,0,0,0,1]

**encode(D, A, G, C) =  $V_D + V_A + V_G + V_C$   
= [1, 0, 1, 0, 0, 0, 1, 0, 0, 0]**

What does it miss?

# One-hot-encoding of $k$ elements

A=[1,0,0,0,0,0,0,0,0,0]

B=[0,1,0,0,0,0,0,0,0,0]

C=[0,0,1,0,0,0,0,0,0,0]

D=[0,0,0,1,0,0,0,0,0,0]

E=[0,0,0,0,1,0,0,0,0,0]

F=[0,0,0,0,0,1,0,0,0,0]

G=[0,0,0,0,0,0,1,0,0,0]

H=[0,0,0,0,0,0,0,1,0,0]

I=[0,0,0,0,0,0,0,0,1,0]

J=[0,0,0,0,0,0,0,0,0,1]

$$\text{encode}(D, A, G, C) = V_D \cdot V_A \cdot V_G \cdot V_C$$

$$=[0,0,0,1,0,0,0,0,0,0, 1,0,0,0,0,0,0,0,0,0, 0,0,0,0,0,0,1,0,0,0, 0,0,1,0,0,0,0,0,0,0]$$

# *Neural* Language Model

***Neural* Language Modeling** runs on a fixed-window, like n-gram:

Input: fixed-window sequence of last  $k$  words:

$$P(x_n \mid x_{n-1}, x_{n-2} \dots x_{n-k}) = \textit{softmax}(\textit{MLP}(x))$$

$$X = \textit{encode}(x_{n-1}, x_{n-2} \dots x_{n-k})$$

# *Neural* Language Model

***Neural* Language Modeling** runs on a fixed-window, like n-gram:

Input: fixed-window sequence of last  $k$  words:

$$P(x_n \mid x_{n-1}, x_{n-2} \dots x_{n-k}) = \textit{softmax}(\textit{MLP}(x))$$

$$X = \textit{encode}(x_{n-1}, x_{n-2} \dots x_{n-k})$$

$$\textit{MLP}(x) = \textit{softmax}(g(g(xW^1 + b^1)W^2 + b^2)W^3 + b^3)$$

# Neural Language Model

**Neural Language Modeling** runs on a fixed-window, like n-gram:

Input: fixed-window sequence of last  $k$  words:

$$P(x_n \mid x_{n-1}, x_{n-2} \dots x_{n-k}) = \text{softmax}(MLP(x))$$

$$X = \text{encode}(x_{n-1}, x_{n-2} \dots x_{n-k})$$

$$MLP(x) = \text{softmax}(g(g(xW^1 + b^1)W^2 + b^2)W^3 + b^3)$$

## Aggregation Option:

$$\begin{aligned} & [0,0,0,1,0,0,0,0,0,0] \\ & \quad + \\ & [1,0,0,0,0,0,0,0,0,0] \\ & \quad + \\ & [0,0,0,0,0,0,1,0,0,0] \\ & \quad + \\ & [0,0,1,0,0,0,0,0,0,0] \\ & \quad = \\ & \mathbf{[1,0,0,1,0,0,1,0,0,0]} \end{aligned}$$

**W**

$$\begin{aligned} A &= [-0.32, 0.09, 0.33, -0.44] \\ B &= [0.29, 0.02, -0.46, -0.39] \\ C &= [-0.46, 0.24, -0.16, 0.08] \\ D &= [-0.15, -0.31, 0.34, 0.00] \\ E &= [-0.10, -0.37, 0.01, 0.40] \\ F &= [-0.28, -0.26, -0.24, 0.31] \\ G &= [-0.32, -0.42, -0.21, 0.18] \\ H &= [-0.09, -0.01, 0.06, 0.14] \\ I &= [0.28, -0.02, -0.39, 0.12] \\ J &= [0.23, -0.22, -0.14, 0.28] \end{aligned}$$

$$(\mathbf{V_D + V_A + V_G + V_C}) \mathbf{W} = \mathbf{V_D W + V_A W + V_G W + V_C W}$$

Sum of rows in W

Each row corresponds to a certain vocabulary item

# Contatination Option:

**W**

A= [-0.32, 0.09, 0.33,-0.44]  
B= [ 0.29, 0.02,-0.46,-0.39]  
C= [-0.46, 0.24,-0.16, 0.08]  
D= [-0.15,-0.31, 0.34, 0.00]  
E= [-0.10,-0.37, 0.01, 0.40]  
F= [-0.28,-0.26,-0.24, 0.31]  
G= [-0.32,-0.42,-0.21, 0.18]  
H= [-0.09,-0.01, 0.06, 0.14]  
I= [ 0.28,-0.02,-0.39, 0.12]  
J= [ 0.23,-0.22,-0.14, 0.28]

$$(V_D \bullet V_A \bullet V_G \bullet V_C) W = ?$$

Contatination Option  
 $(V_D \bullet V_A \bullet V_G \bullet V_C) W = ?$

$[0,0,0,1,0,0,0,0,0,0]$

0

$[1,0,0,0,0,0,0,0,0,0]$

0

$[0,0,0,0,0,0,1,0,0,0]$

0

$[0,0,1,0,0,0,0,0,0,0]$

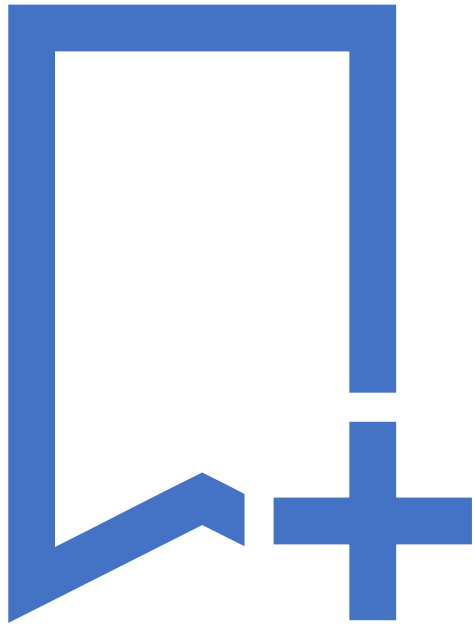
=

**$[0,0,0,1,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,0,0,0,1,0,0,0,0,0,1,0,0,0,0,0,0,0]$**

still sum of rows in W but W has 4x many rows

D(-3)= [-0.12,-0.24, 0.12,-0.34]  
E(-3)= [-0.42,-0.21, 0.08, 0.40]  
F(-3)= [ 0.20, 0.11,-0.31, 0.33]  
G(-3)= [ 0.07,-0.05, 0.16, 0.23]  
H(-3)= [ 0.28, 0.03, 0.22,-0.49]  
I(-3)= [ 0.08, 0.39,-0.25, 0.27]  
J(-3)= [ 0.10,-0.42,-0.37, 0.35]  
A(-2)= [-0.00, 0.41, 0.19, 0.49]  
B(-2)= [ 0.24, 0.48, 0.34,-0.42]  
C(-2)= [-0.46, 0.22, 0.24,-0.21]  
D(-2)= [-0.11,-0.48, 0.18,-0.22]  
E(-2)= [-0.32, 0.10,-0.41,-0.43]  
F(-2)= [ 0.32, 0.02,-0.22, 0.06]  
G(-2)= [-0.31,-0.36, 0.09, 0.39]  
H(-2)= [ 0.01,-0.22,-0.09,-0.15]  
I(-2)= [ 0.01, 0.10,-0.16,-0.21]  
J(-2)= [-0.24, 0.40,-0.34,-0.13]  
A(-1)= [-0.23,-0.38, 0.02, 0.32]  
B(-1)= [-0.34, 0.04,-0.18,-0.00]  
C(-1)= [ 0.40,-0.02, 0.10,-0.16]  
D(-1)= [ 0.13,-0.07,-0.19,-0.01]  
E(-1)= [ 0.40, 0.27,-0.33, 0.36]  
F(-1)= [ 0.04,-0.13,-0.43, 0.39]  
G(-1)= [ 0.44, 0.38, 0.03,-0.39]  
H(-1)= [ 0.41,-0.23, 0.33,-0.08]  
I(-1)= [-0.50,-0.16,-0.42,-0.27]  
J(-1)= [-0.15, 0.41, 0.46,-0.16]  
A(+0)= [-0.11, 0.03, 0.20, 0.50]  
B(+0)= [ 0.16,-0.34, 0.20,-0.21]  
C(+0)= [ 0.05,-0.13,-0.23,-0.31]





- 1-hot times matrix: matrix row selection
- Sum of 1-hot times matrix: row selection + sum
- Concat of 1-hot: like using 1-hot from larger vocab

# Embedding Layer

**E**

A= [-0.32, 0.09, 0.33,-0.44]  
B= [ 0.29, 0.02,-0.46,-0.39]  
C= [-0.46, 0.24,-0.16, 0.08]  
D= [-0.15,-0.31, 0.34, 0.00]  
E= [-0.10,-0.37, 0.01, 0.40]  
F= [-0.28,-0.26,-0.24, 0.31]  
G= [-0.32,-0.42,-0.21, 0.18]  
H= [-0.09,-0.01, 0.06, 0.14]  
I= [ 0.28,-0.02,-0.39, 0.12]  
J= [ 0.23,-0.22,-0.14, 0.28]

$\text{encode}(D, A, G, C)$

$$= \mathbf{E}_{[D]} \circ \mathbf{E}_{[A]} \circ \mathbf{E}_{[G]} \circ \mathbf{E}_{[C]}$$

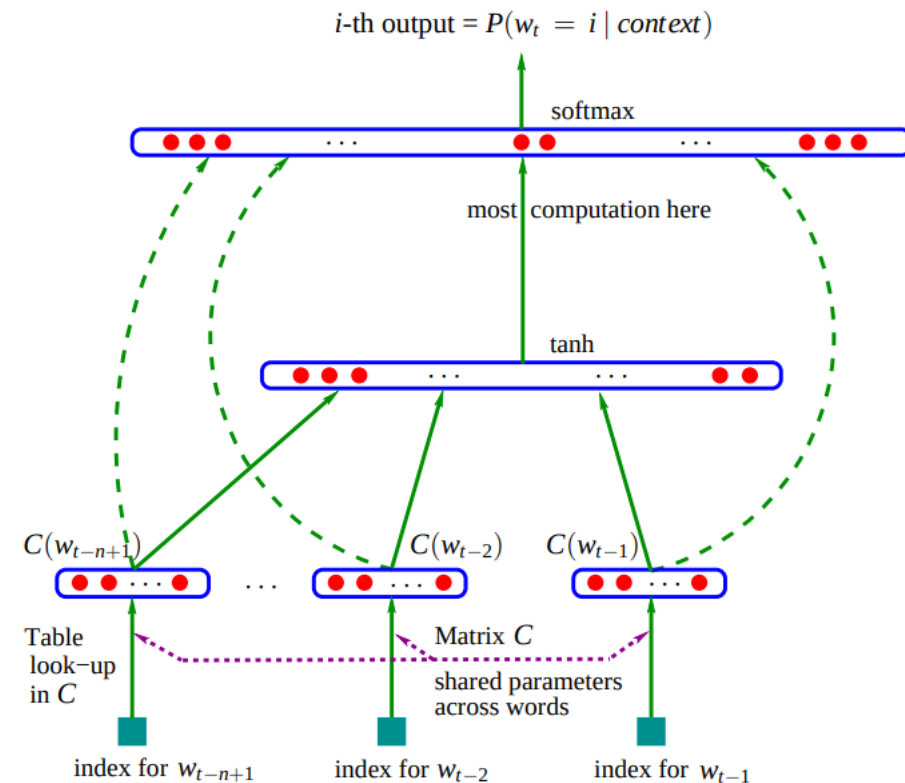
[-0.15,-0.31, 0.34, 0.00,-0.32, 0.09, 0.33,-0.44,-0.32,-0.42,-0.21, 0.18,-0.46, 0.24,-0.16, 0.08]

- Assigns each item of the vocabulary a unique number
- Associate it with a row in matrix **E** of dense vectors (row dimension  $\ll |V|$  )
- concat or sum rows of **E** for the given input

# A Neural Probabilistic Language Model

- Bengio et al. (2003)

BENGIO, DUCHARME, VINCENT AND JAUVIN



1: Neural architecture:  $f(i, w_{t-1}, \dots, w_{t-n+1}) = g(i, C(w_{t-1}), \dots, C(w_{t-n+1}))$  where neural network and  $C(i)$  is the  $i$ -th word feature vector.



## Training a neural language model

- Set dimensions of the layers  $E$ ,  $W_3$ , according to your vocab size.
- Initialize with random values for  $E$ ,  $W_1$ ,  $W_2$ ,  $W_3$ ,  $b_1$ ,  $b_2$ ,  $b_3$
- For every  $n$ -tuple in some text:
  - try to predict last item based on prev  $n-1$
  - use cross-entropy loss

# Neural Language Model can do:



Probability score for a given sentence



Generate new sentences



Predict next word  $i$  based on previous  $k$  words



Predict word label based on  $k$  items (when?)

# What happens after the training?



Consider the columns of the last layer  $W_3$ .



Consider the rows of the embedding layer  $E$ .



The columns of  $W_3$  corresponds to the vocabulary items (!)



# Review

Note the most important thing you've learnt so far

# The issue with one-hot-vector

```
motel = [0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]  
hotel = [0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]
```





# Issues with the previous method

```
motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
hotel = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]
```

- Words are treated as discrete symbols: no sense of similarity. The vectors are orthogonal
- Vector Dimension = # of words in the vocabulary.
- Training a language model is expensive (why?)
- We want **better vector representation**

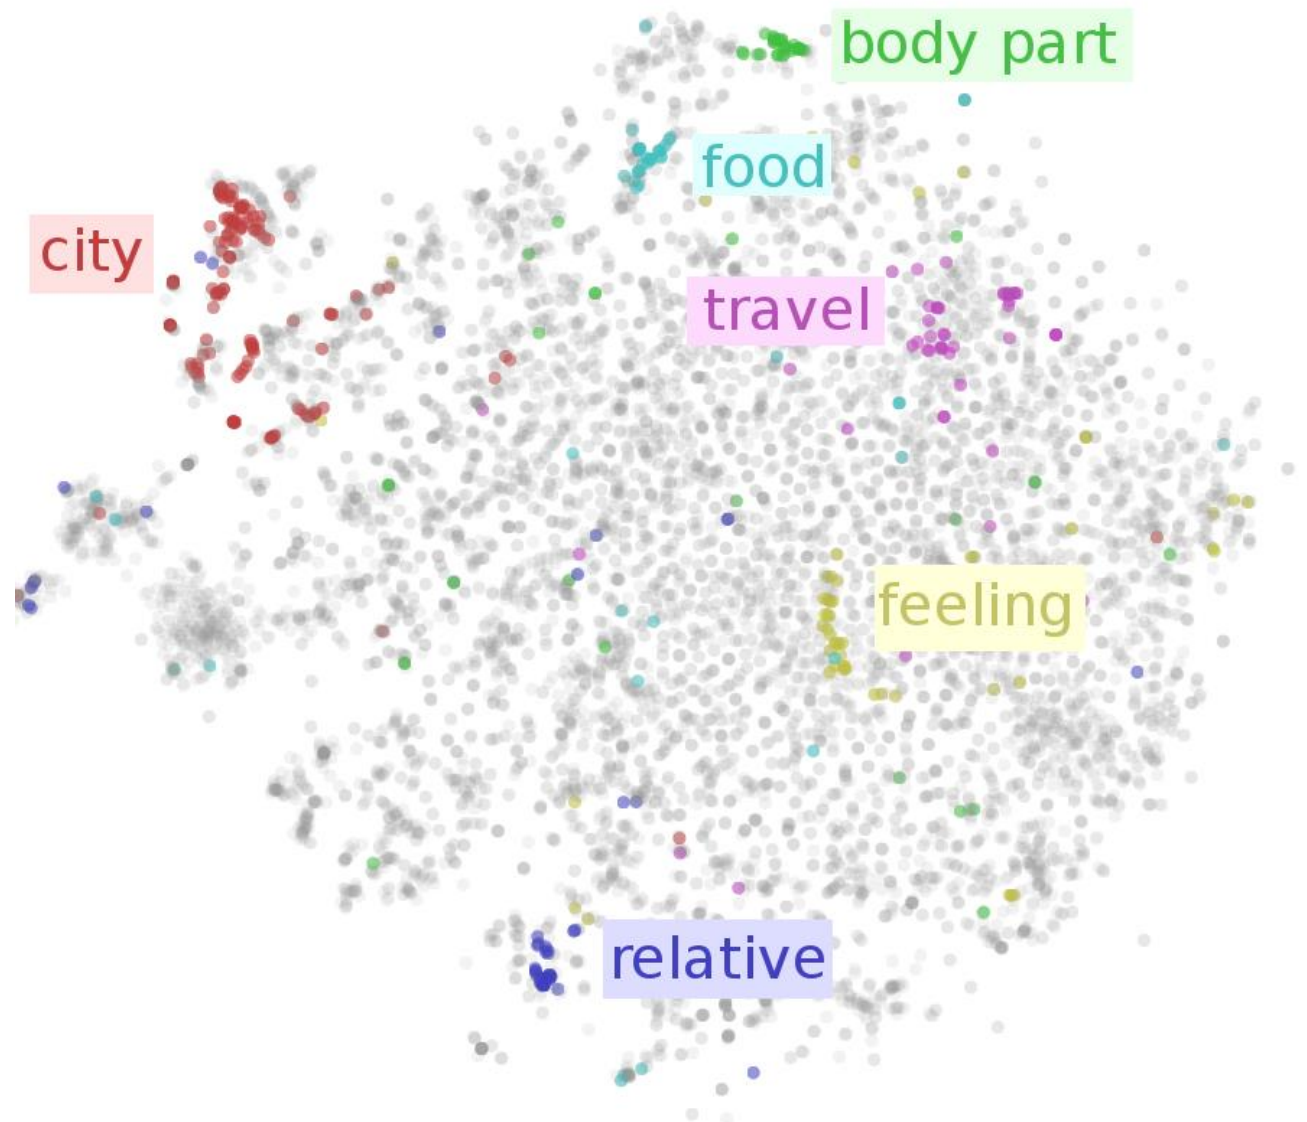
A large, solid orange oval shape that serves as the background for the text.

How?

---

# Word Vectors

- Aka:  
*Word Embedding /*  
*Word Representations /*  
*Distributed Representation*



Your Turn!

Word2Vec  
Family  
Algorithm

You have **1h** to prepare short presentation about **Distributional Semantics**:

- *Group A*: Word2Vec
- *Group C*: Neural Word Embedding as Implicit Matrix Factorization
- *Group D*: gloVe
- *Group E*: FastText

# Prepare to teach the rest:

- Read The Papers
- Read additional Material
- Prepare a presentation (~20 min long)
- Points to Cover:
  - Model structure
  - Training process
  - Output vectors semantic properties
  - Differences (l.e., from the language model)
  - Critics



# Distributional Semantics



A word's meaning is given by the words that frequently appear with it.



“You shall know a word by the company it keeps”  
(J. R. Firth 1957: 11).



When a word  $w$  appears in a text, its context is the set of words that appear nearby (within a fixed-size window).



Neural Language Model can use this context to build a representation of  $w$ .

# A fixed-window neural Language Model

We need a neural network  
that can operate on  
*variable lengths* of input

## Improvements over n-gram LM:

- No sparsity problem
- Don't need to store all observed n-grams

## Remaining problems:

- Fixed window is too small
- Enlarging window makes  $W$  bigger
- Window can never be large enough!
- Every word is multiplied by **completely different weights** in  $W$ .  
No symmetry in how the inputs are processed.

# Language Models – Recommended Reading

- <https://lilianweng.github.io/lil-log/2019/01/31/generalized-language-models.html>