Recurrent Neural Networks

Liad Magen

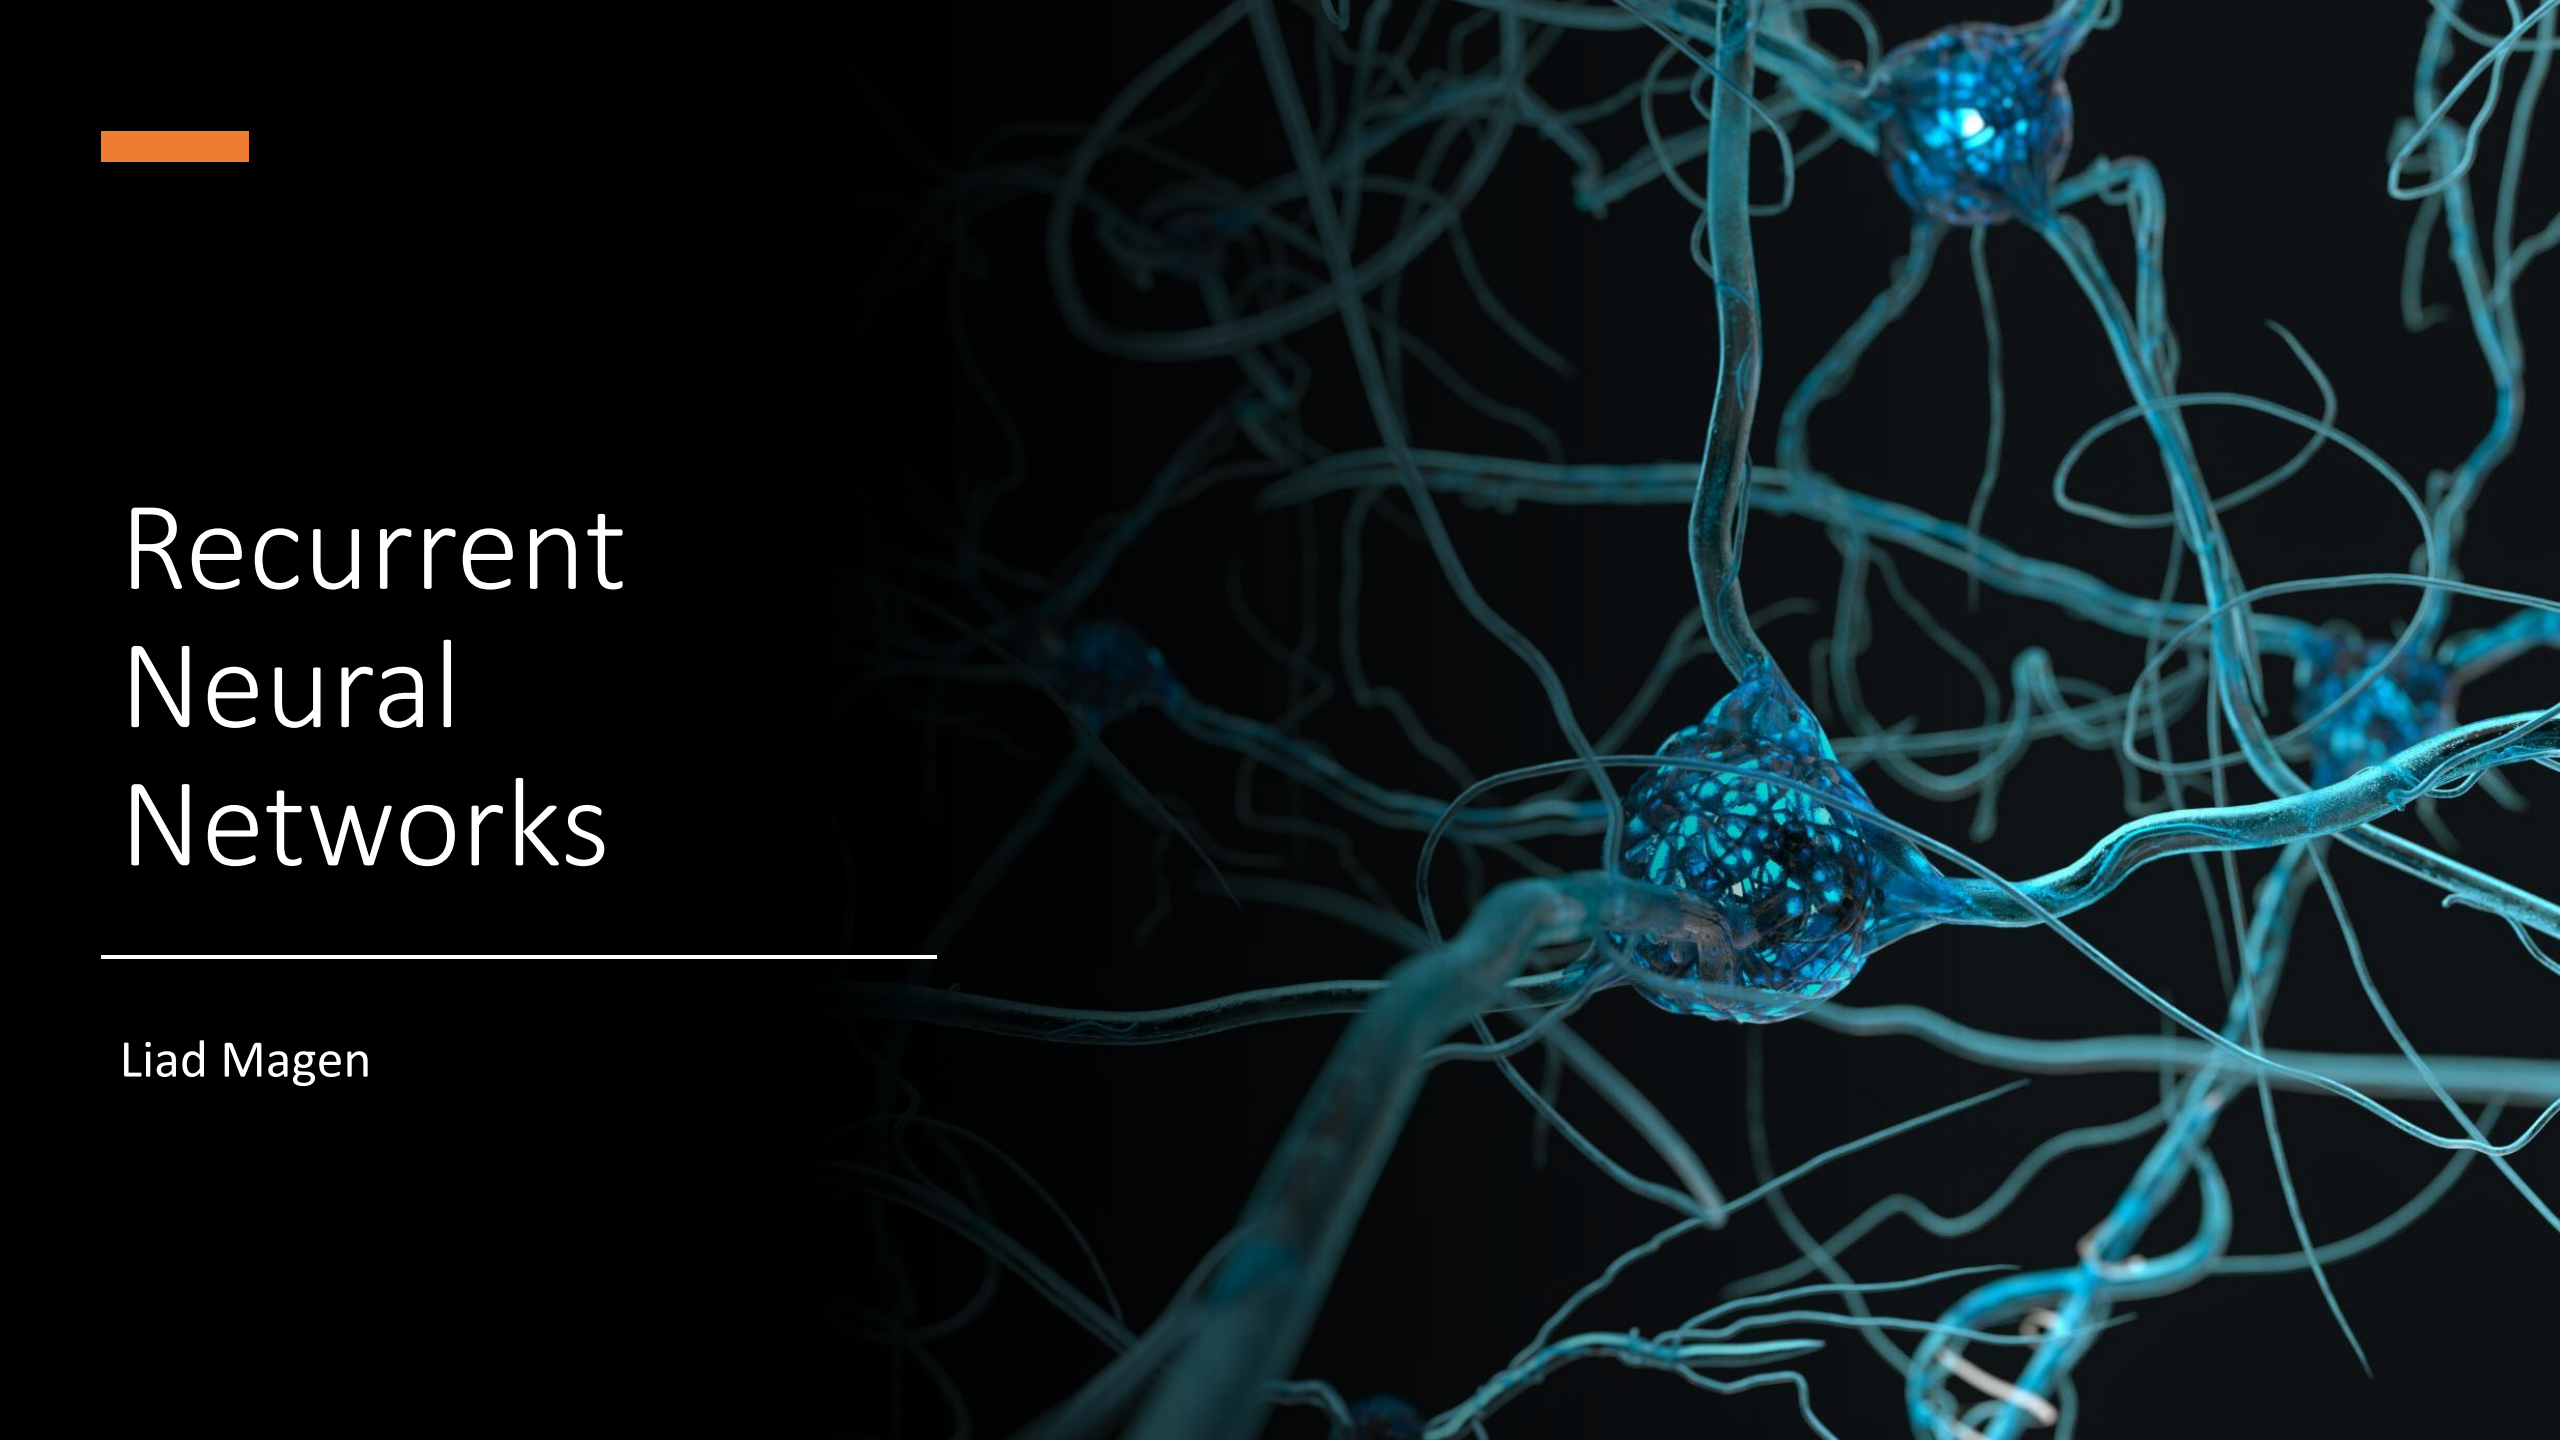

Long-Distance dependencies in Language

The problem of a fixed-window Neural Language Model:

The trophy would not fit in the brown suitcase because *it* is too **big**.

Trophy

The trophy would not fit in the brown suitcase because *it* is too **small**.

Suitcase

Commonsense Reasoning – Winograd Schema Challenge

Recurrent Neural Network (RNN) (Elman 1990)

- An unrolled RNN:
  - Very deep Feed-Forward Network
  - Shared parameters across the layers
  - Can get a *new input* at each layer

From NN to RNN

RNN feeds the output of a previous layer together with the next input to predict the output of the layer.

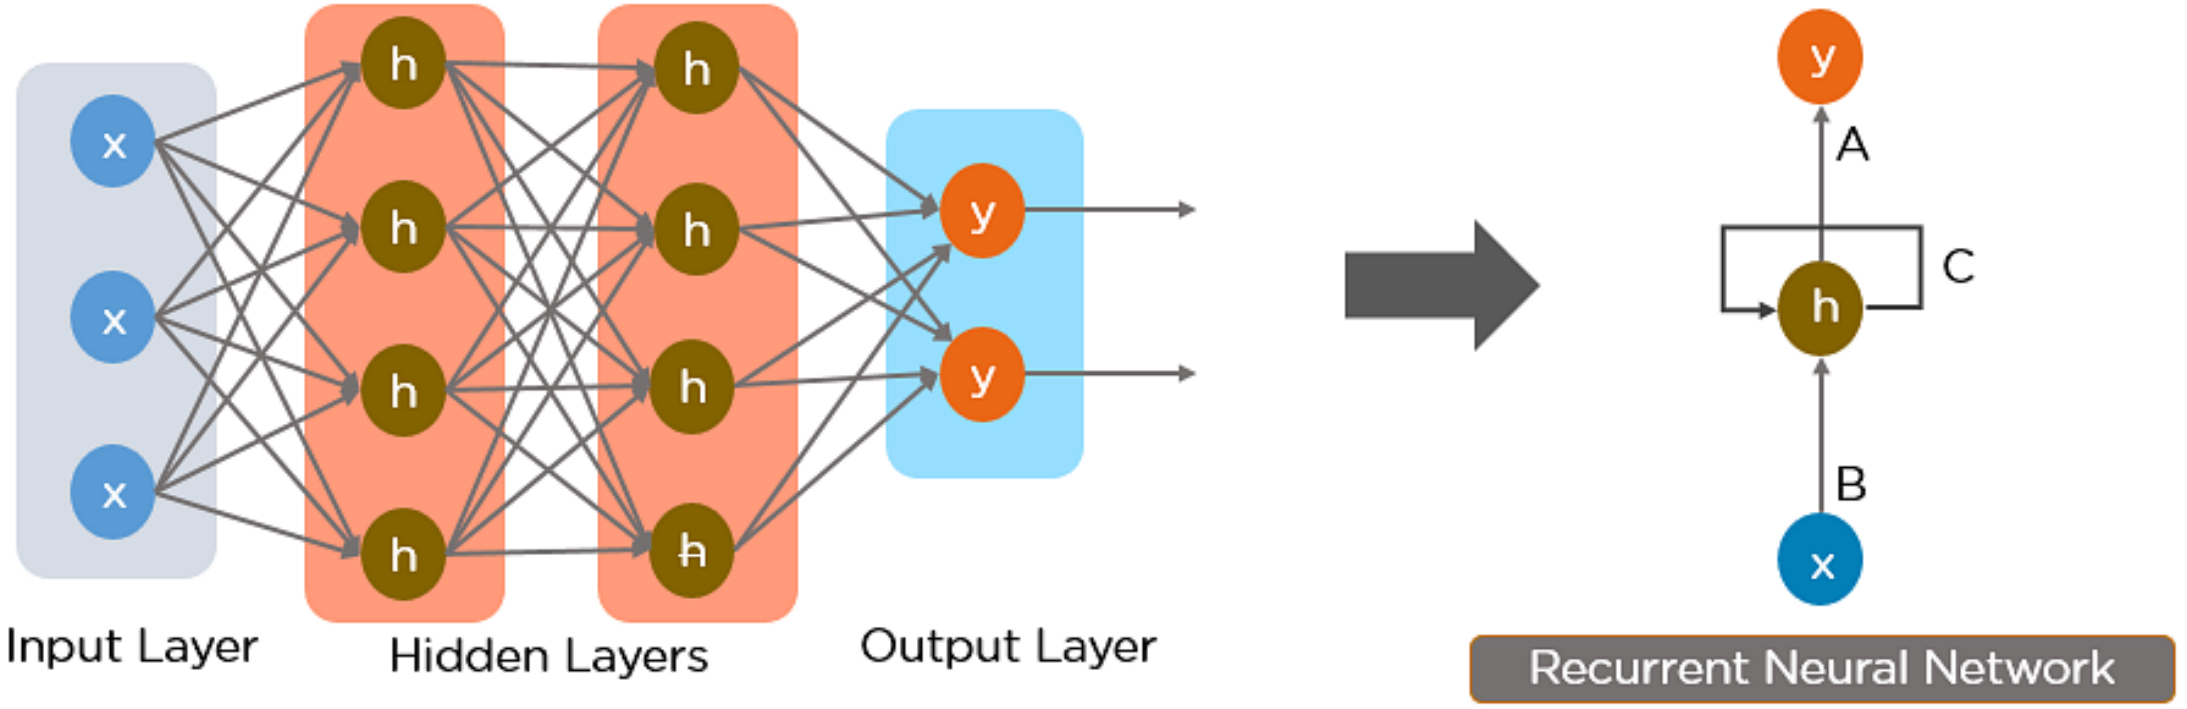

RNN – information cycles through the hidden layer:

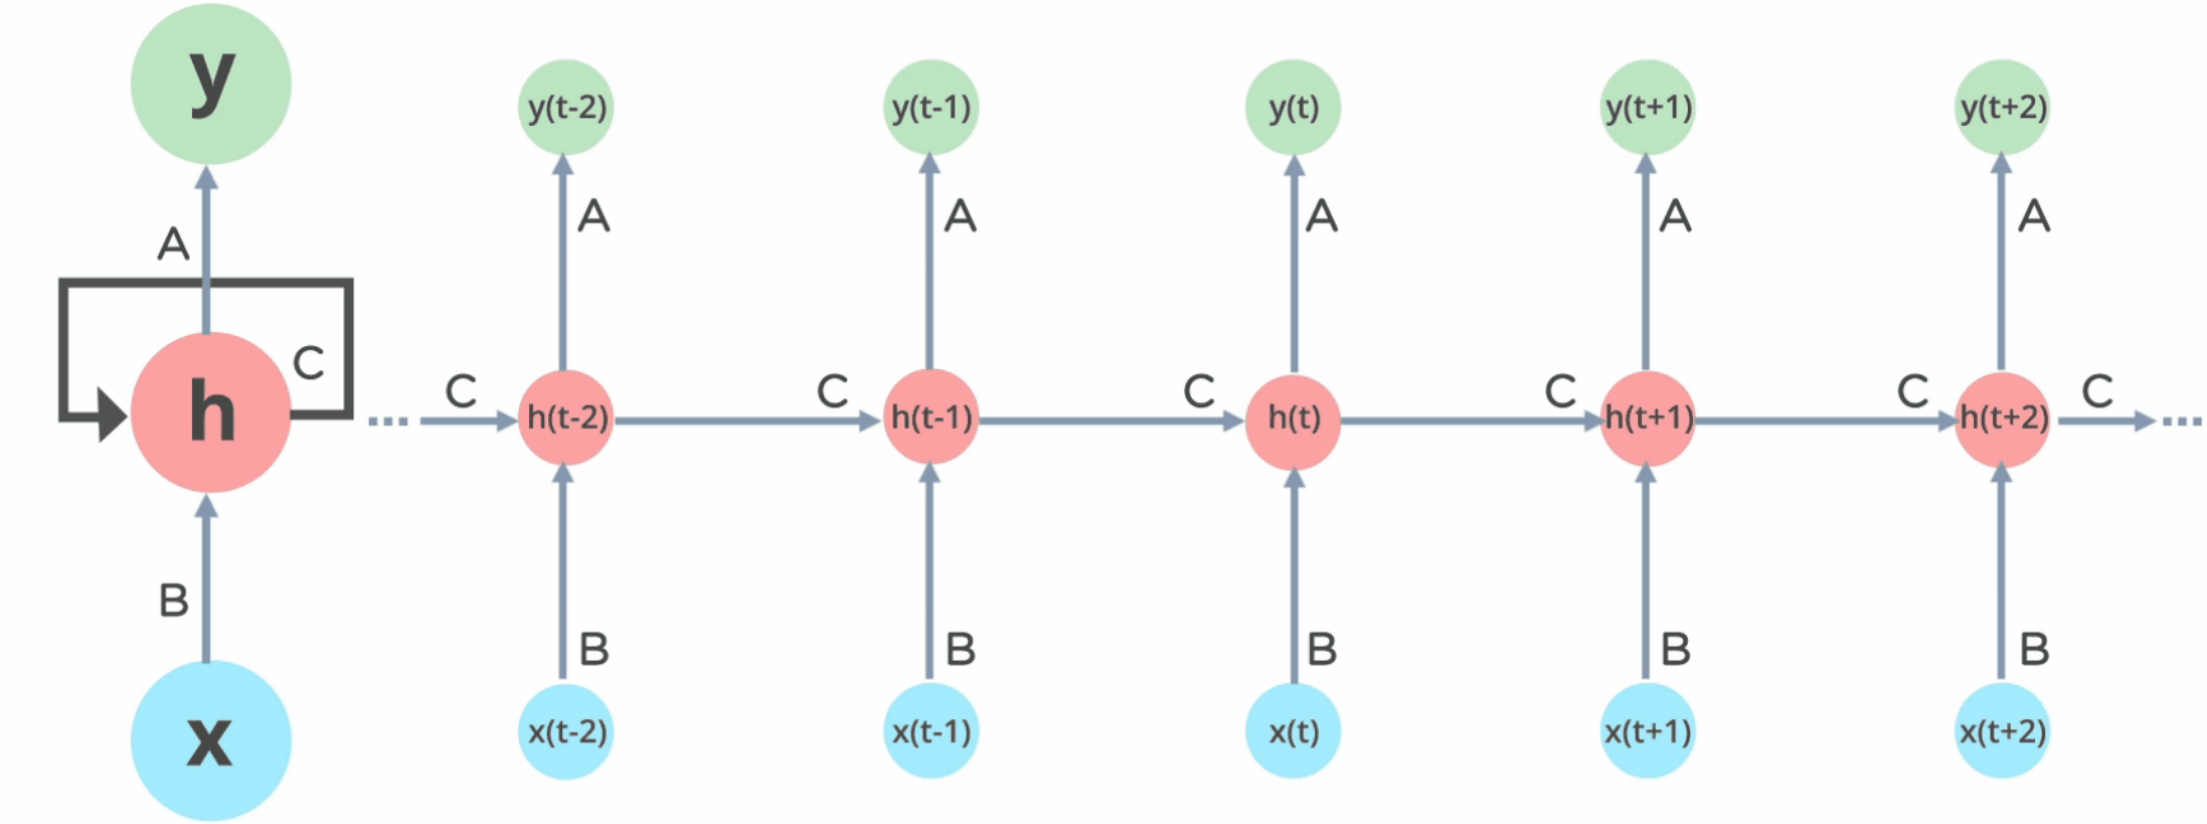

Example – What time is it ?

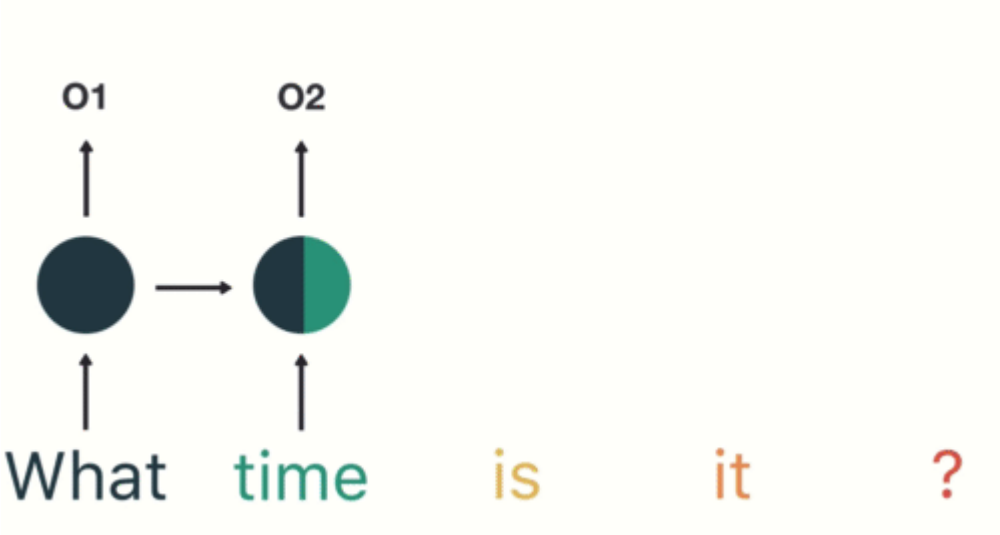

Recurrent Neural Networks

- Very strong models of sequential data
- **Trainable** function (n vectors) --> single vector
- Input: a set of word-vectors: v1, v2, v3...
- Apply the same weight – $\mathbb{W}$ - repeatedly
- What is the output?

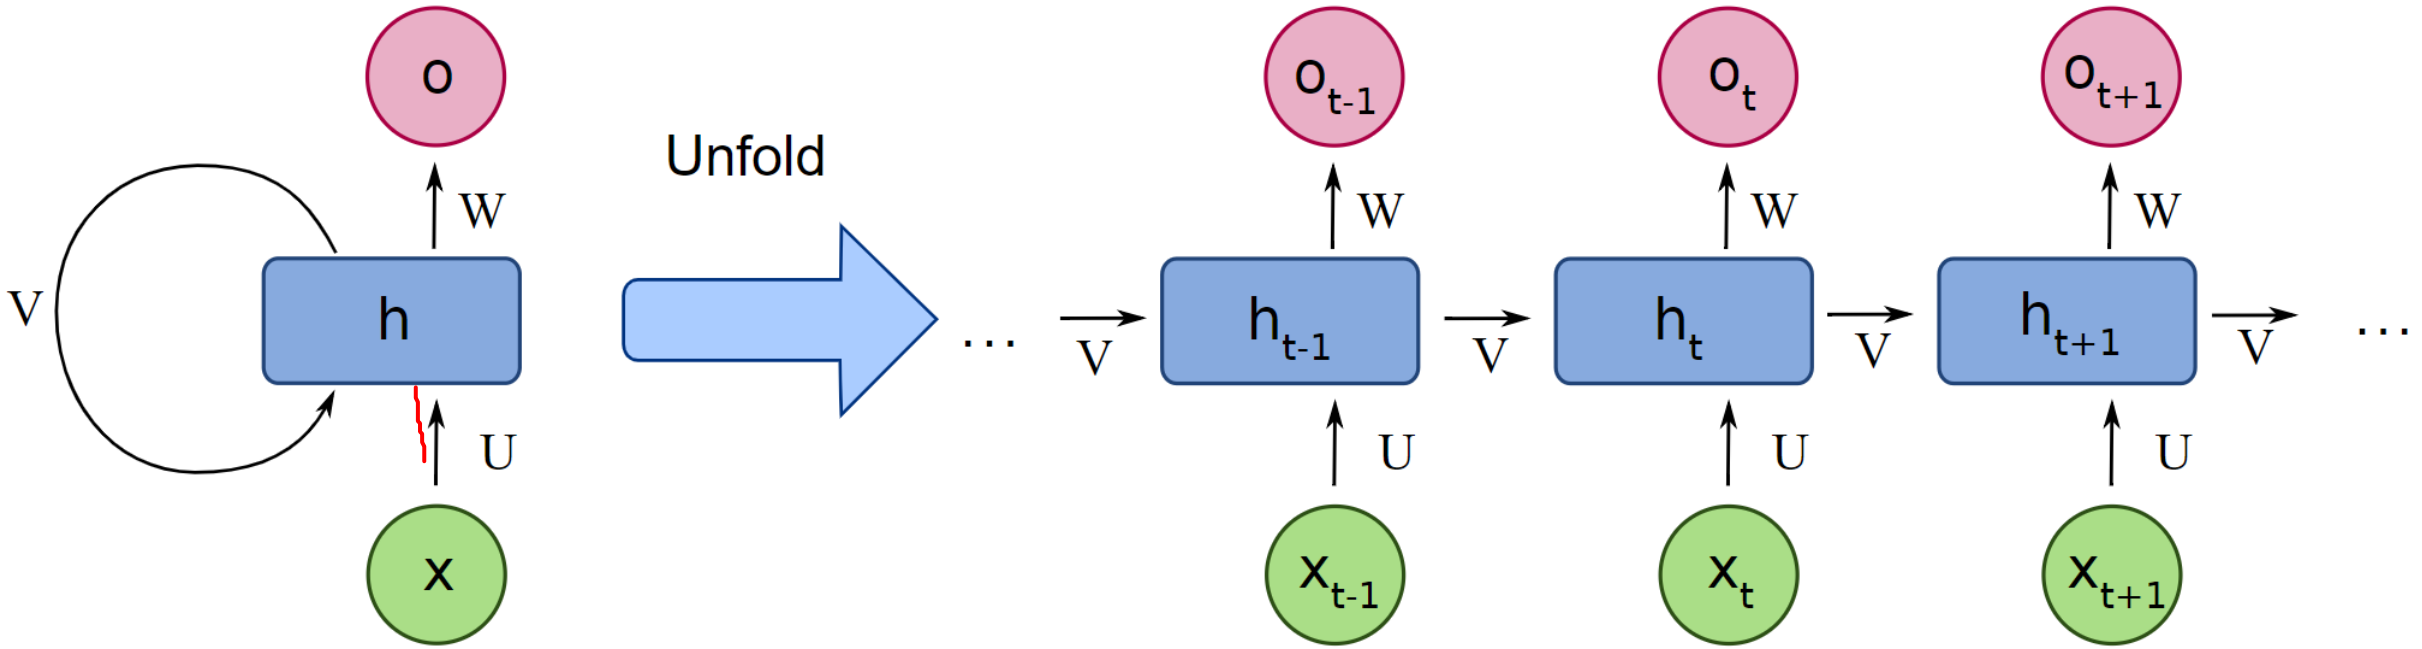

RNN Intermediate Output Vectors

What are the output vectors (O_t) good for?

- By default? For Nothing.
  But we can train them!
  - Define function form
  - Define loss

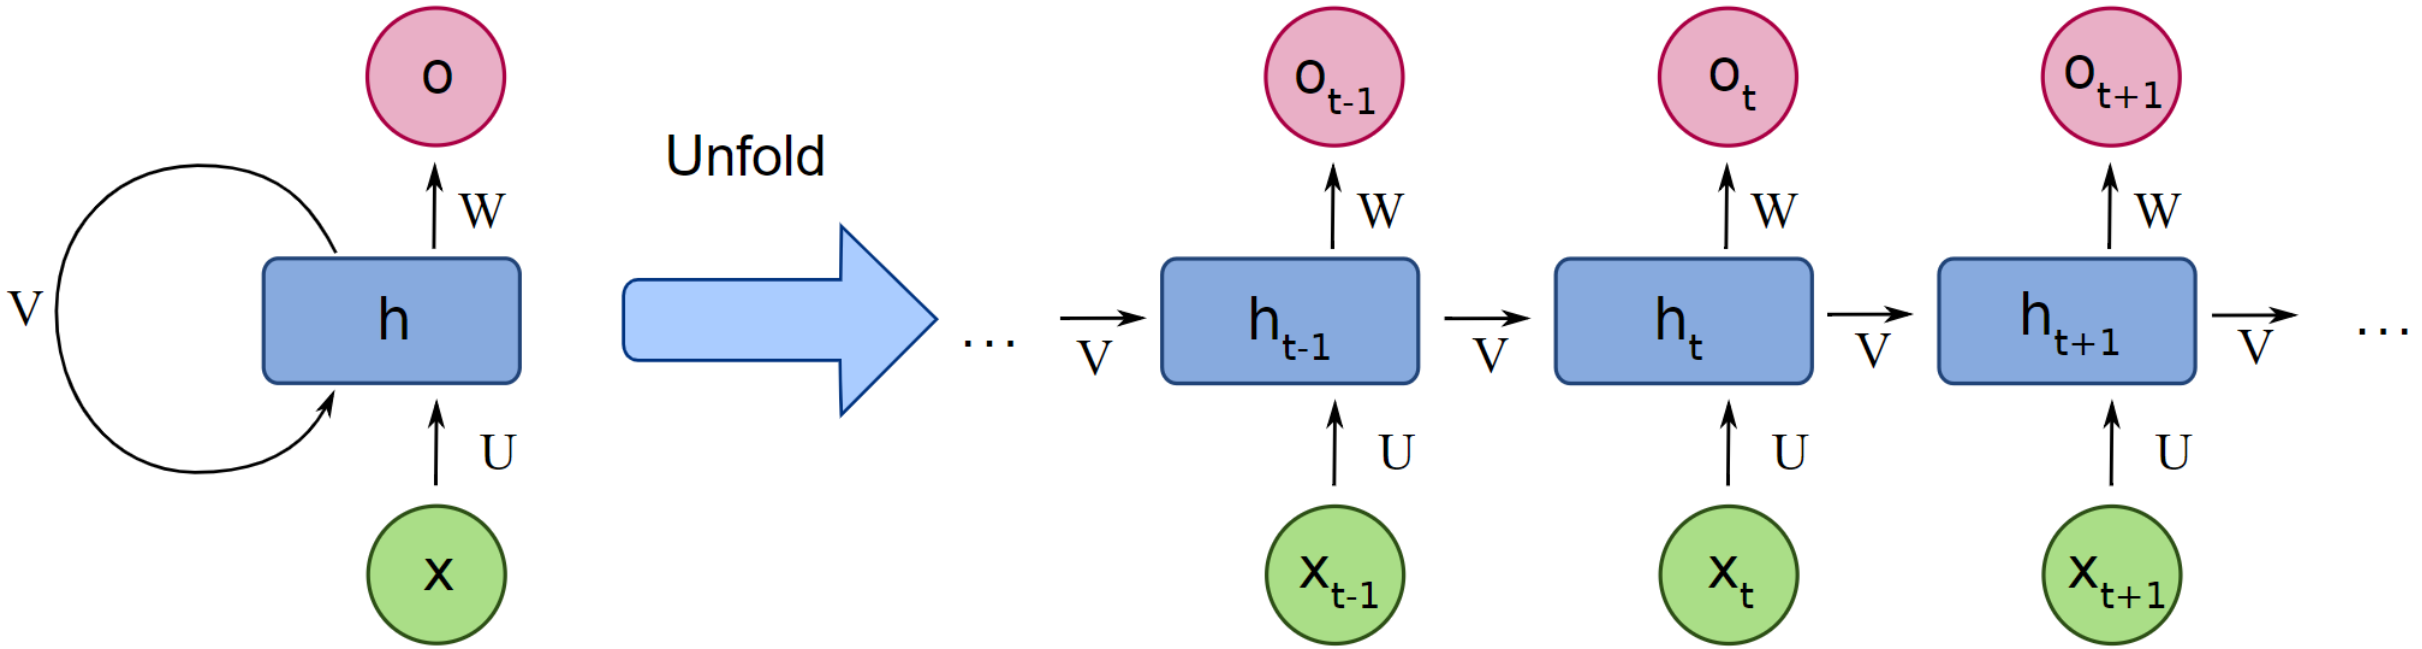

RNN Intermediate Output Vectors

What are the output vectors (O_t) good for
Examples:

- Next word in the sequence
- Recommednation: Next item to buy/watch
- Tagging: POS, NER, SRL

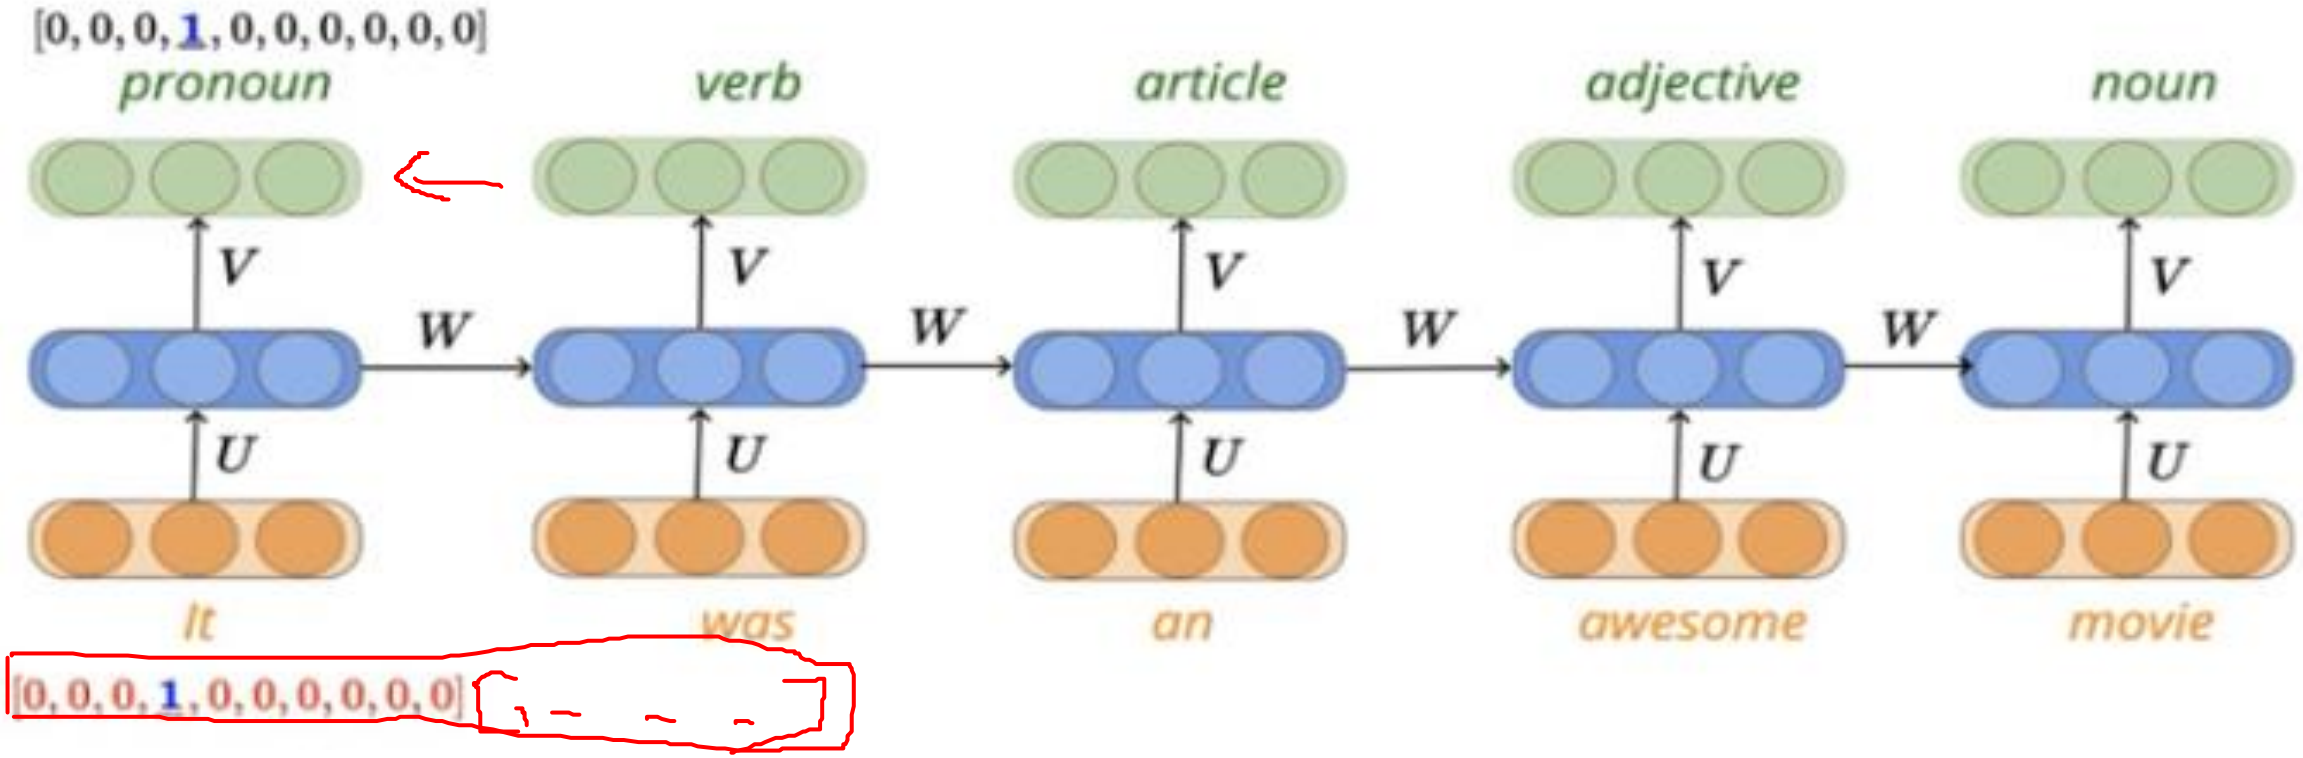

What can RNN do?

- Represent a sentence as a vector
  - Read a whole sentence, make a prediction
- Represent a context in a sentence
  - Read context up until a time-point.
- Generate content

Content Generation

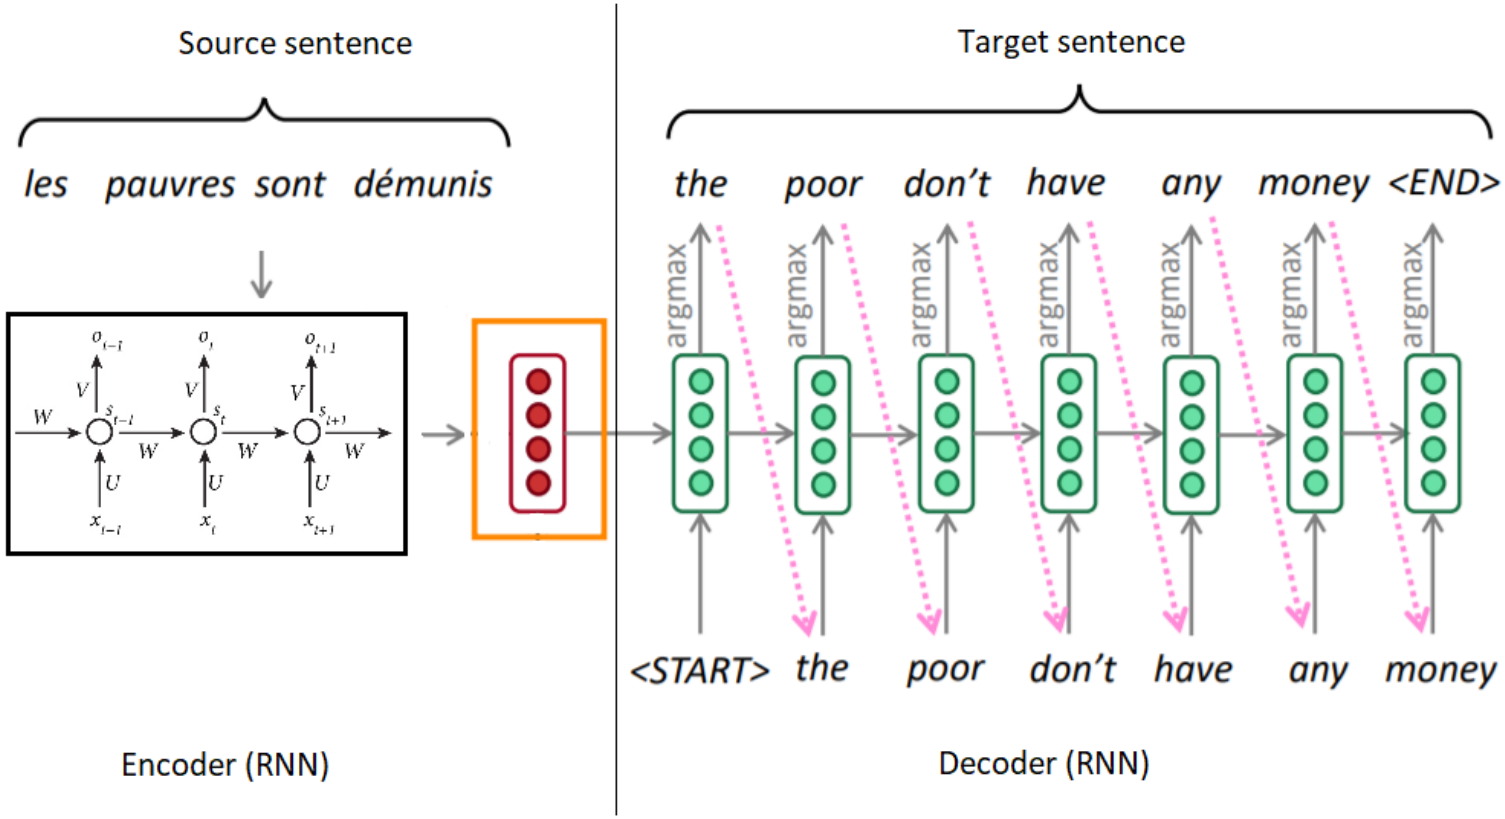

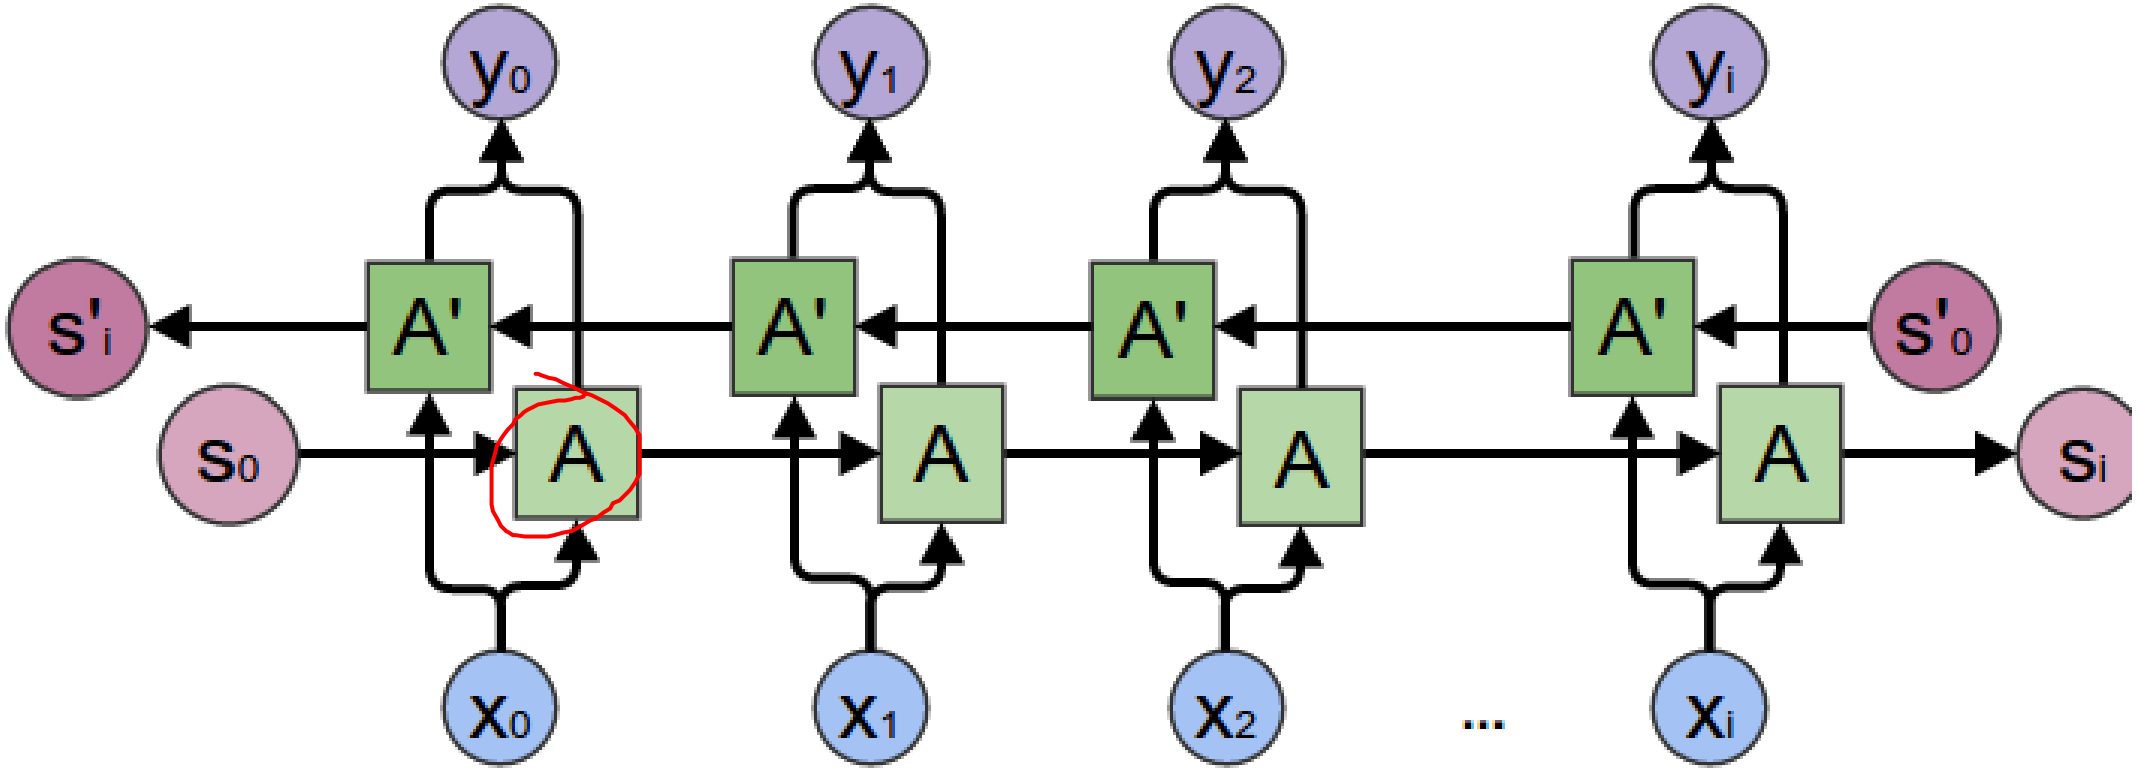

Bidirectional RNNs

One RNN runs left to right.

Another runs right to left.

Encode both future and history of a word.

(An infinite window, around the word)

Deep Bi-RNNs

- The hidden layer can be stacked.

- Provides an 'infinite' deep window around a focus word
- Learn to extract what's important
- Easy to train
- Very effective for sequence tagging

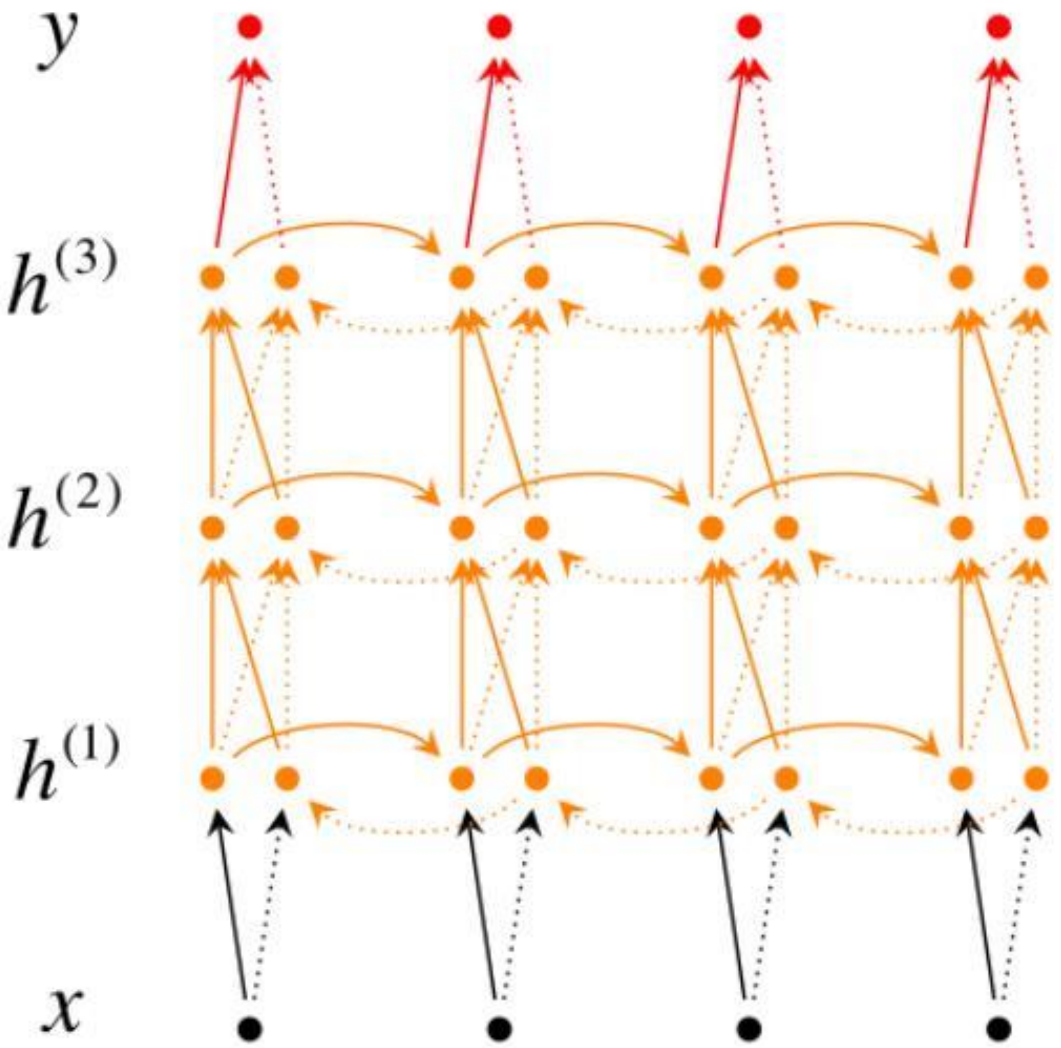

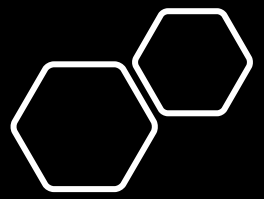

Attention Mechanism (2015)

- “You can’t cram the meaning of a whole $@#ing sentence into a single $@#ing vector”.
  - Ray Mooney

- Solution:
  - Each word is encoded into a vector
  - When decoding, perform a weighted linear combination of these vectors
  - Use the combination in picking the next word

Attention weights

Convolution

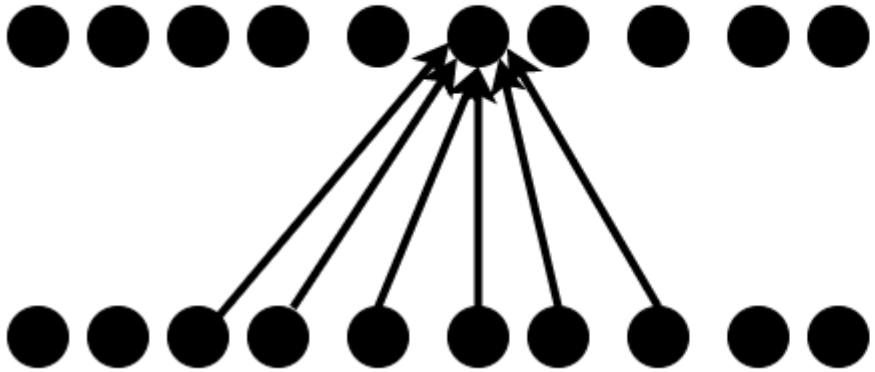

Global attention

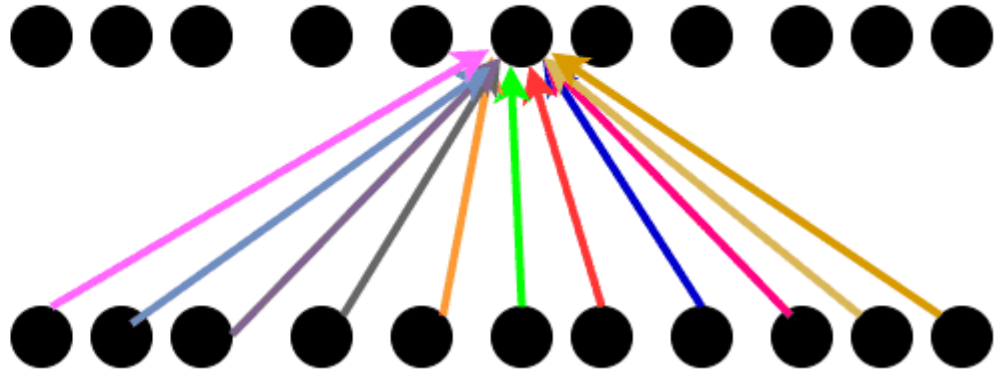

Fully Connected layer

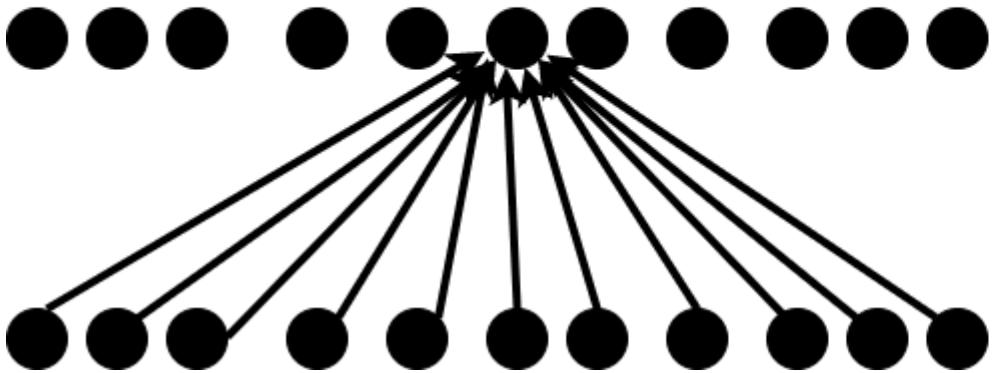

Local attention

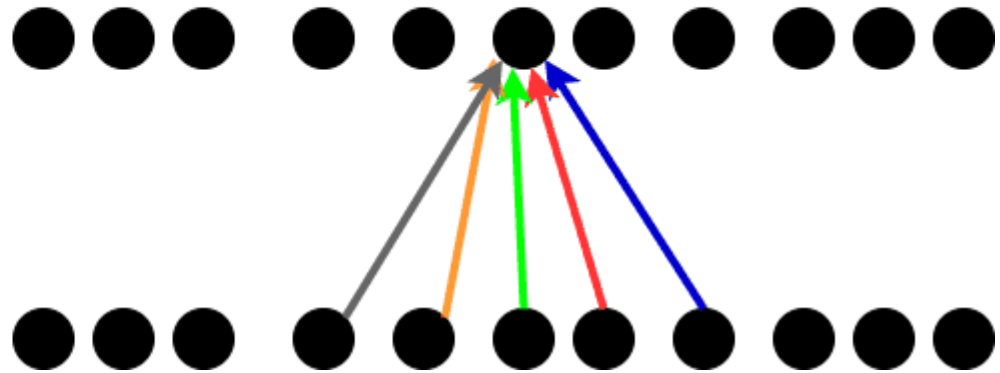

Attention

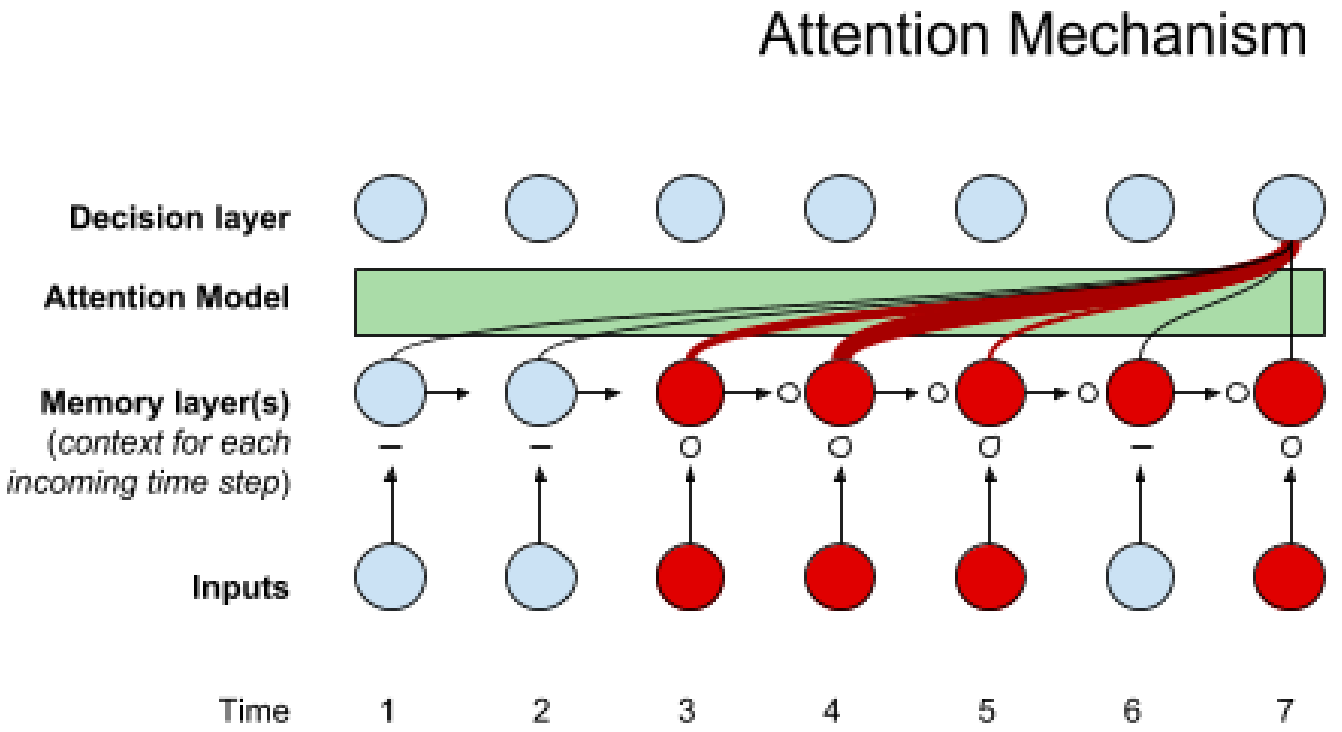
Attention Mechanism
Decision layer
Attention Model
Memory layer(s)
(context for each
incoming time step)
Inputs
Time
1
2
3
4
5
6
7

Ideas to try it out:

Train on Game of thrones to write the last book...

... Or a new Harry Potter:

Test yourself

- What is a Language Model?

A system that predicts the next word

- How is the Linguistic field of human sounds called?

Phonetics / Phonology

- How is the field of words and their parts called?

Morphology

BI-RNN's Recap

- Represent the history up to a point in the sequence, and the future from a point in a sequence.
- Feed into an MLP (or linear classifier) to classify the point based on history and future.
- The network learns which features are important in the history and future for the given prediction task.
- "Infinite window"

Multi-task Learning

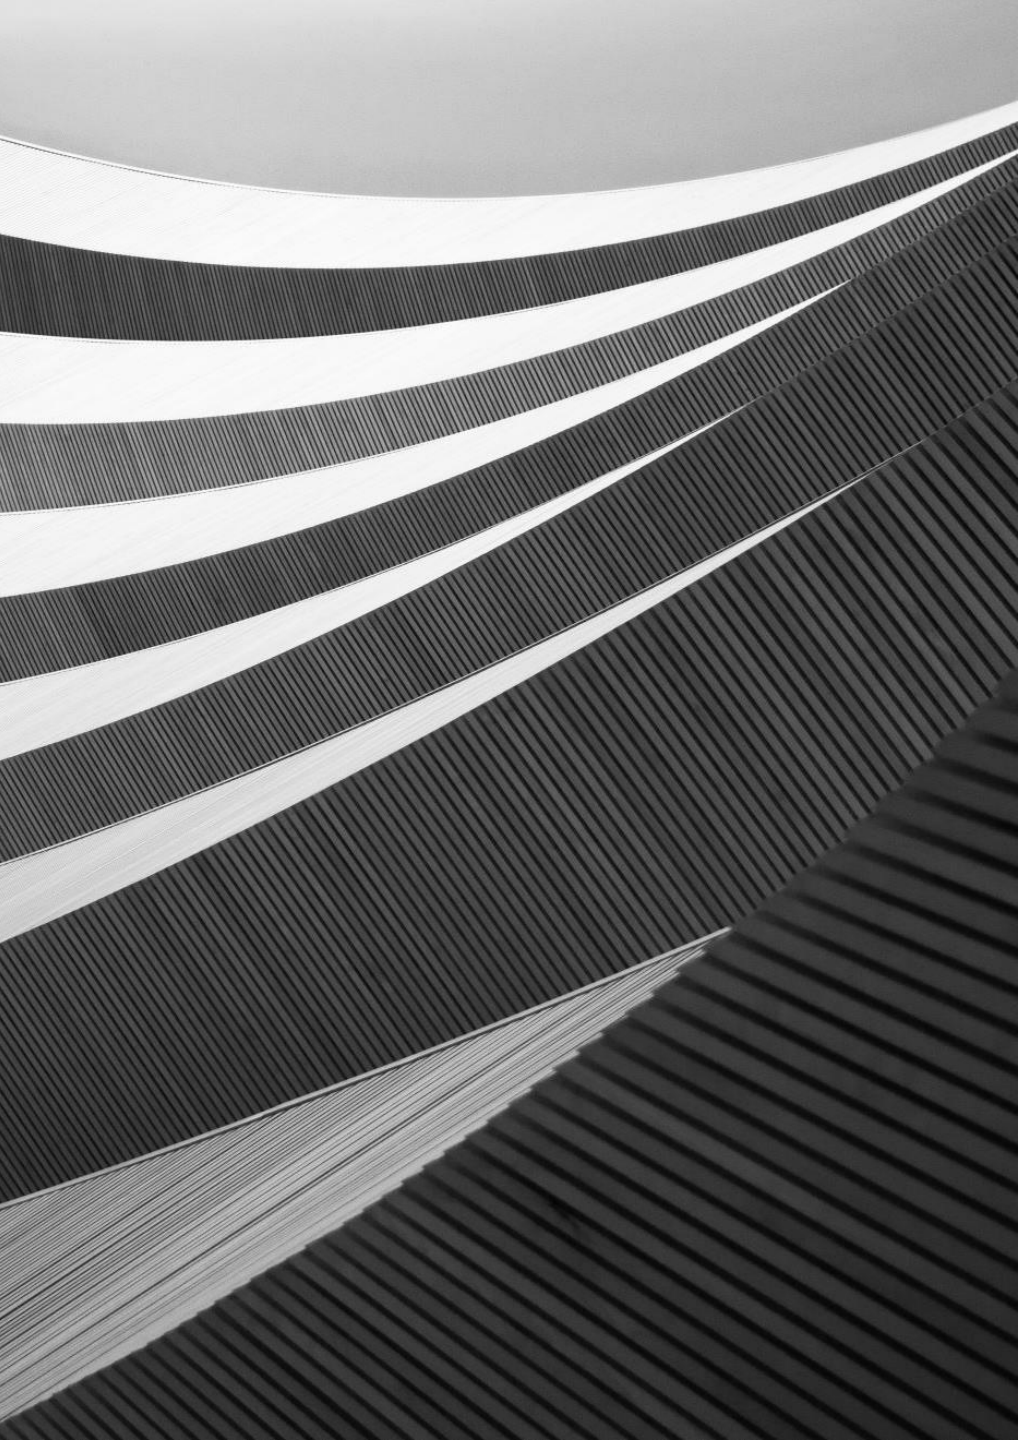

Multi-task Learning

Different sequence prediction tasks have shared structures.

Hints for predicting A may help to predict B.

Instead of training a network to do one thing, train it to do several things.

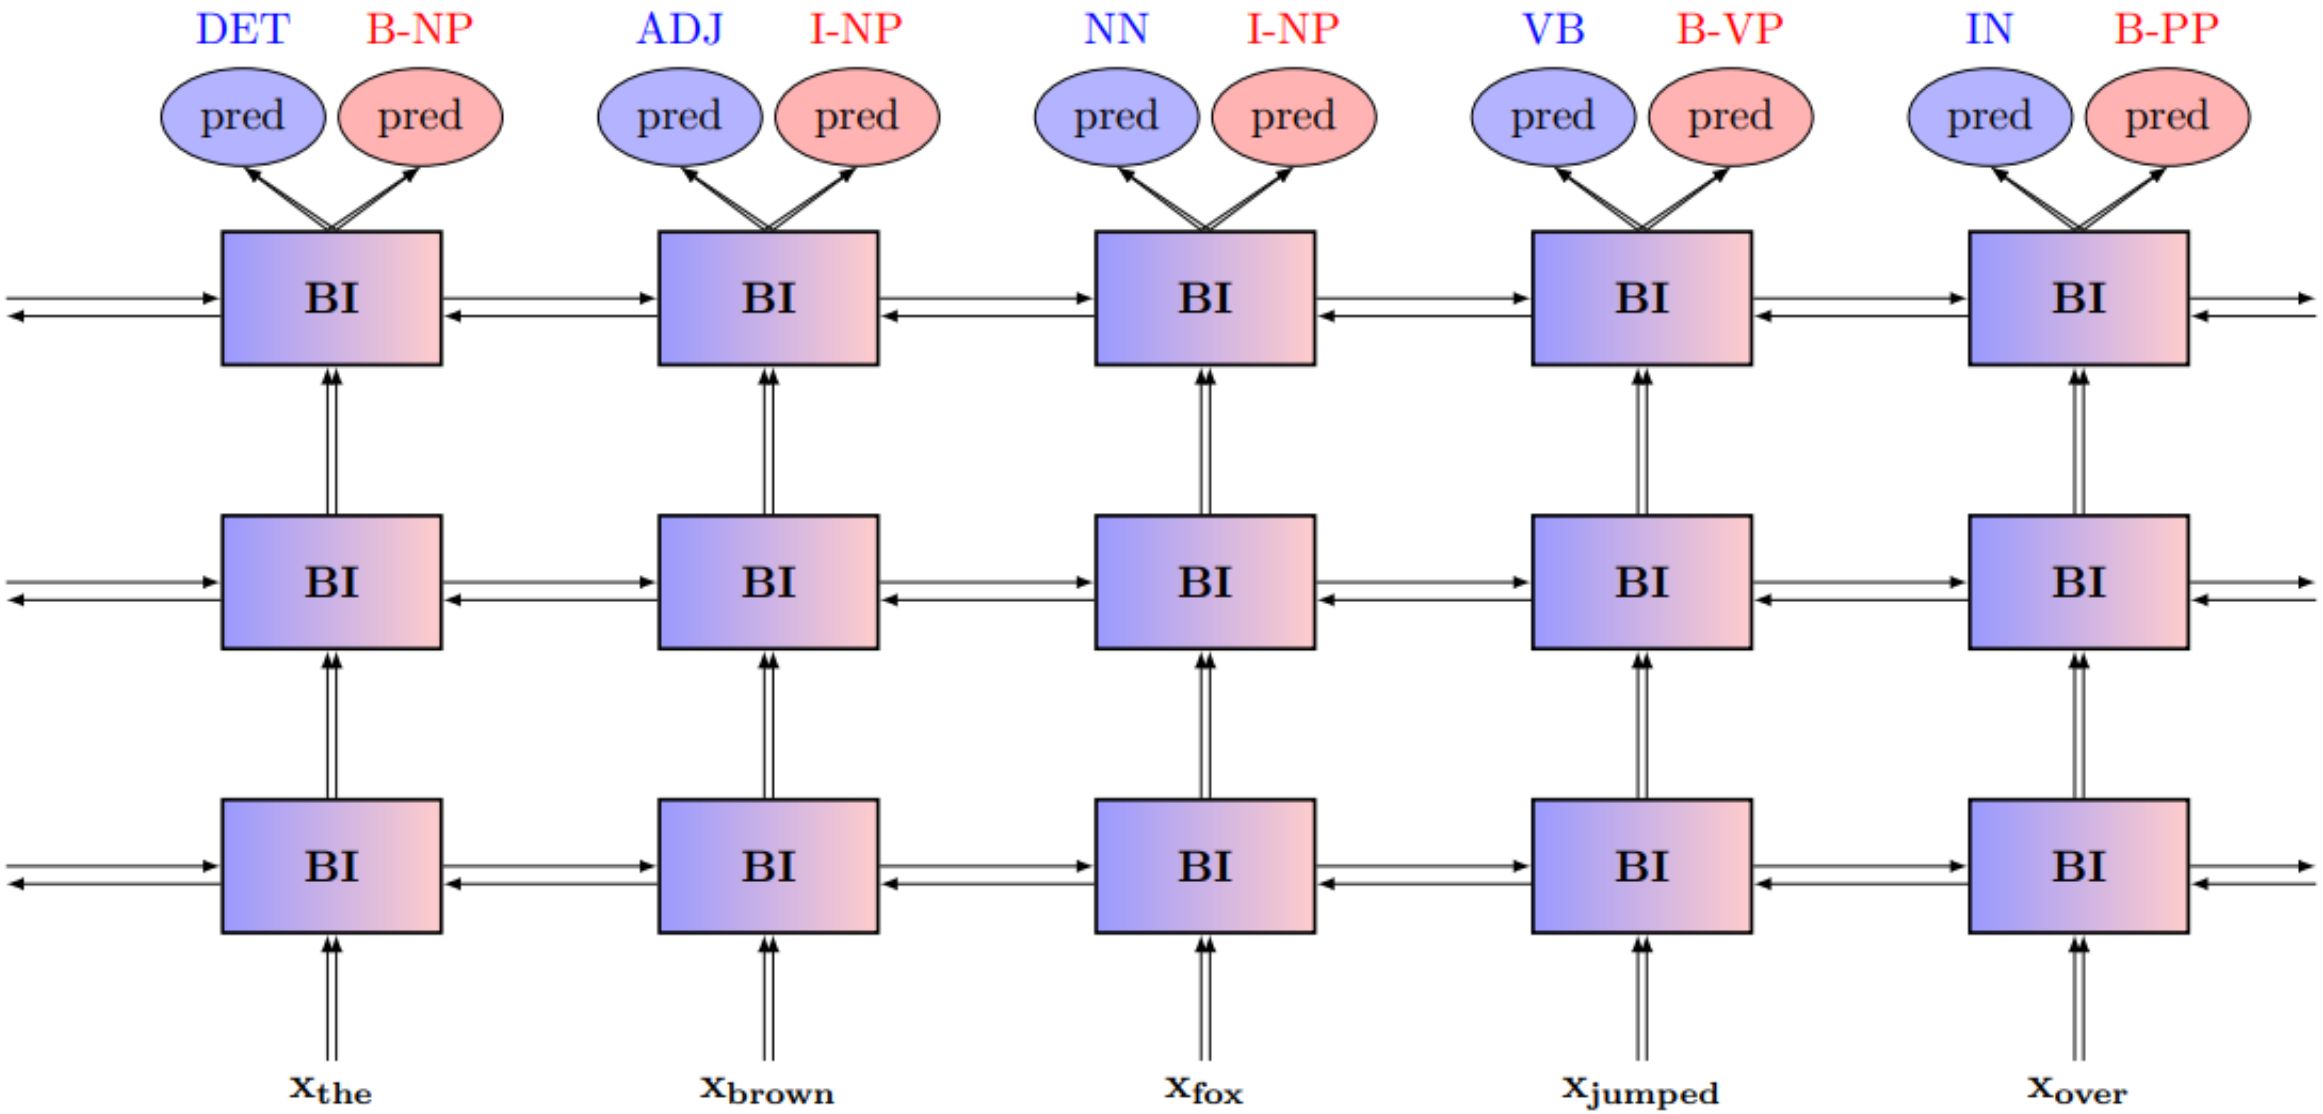
DET
B-NP
ADJ
I-NP
NN
I-NP
VB
B-VP
IN
B-PP
pred
BI
x_{the}
x_{brown}
x_{fox}
x_{jumped}
x_{over}

But...

Not easy to get it to work

For many task pairs – there's no improvement, at all

If the network not wide/deep enough, MTL hurts both tasks...

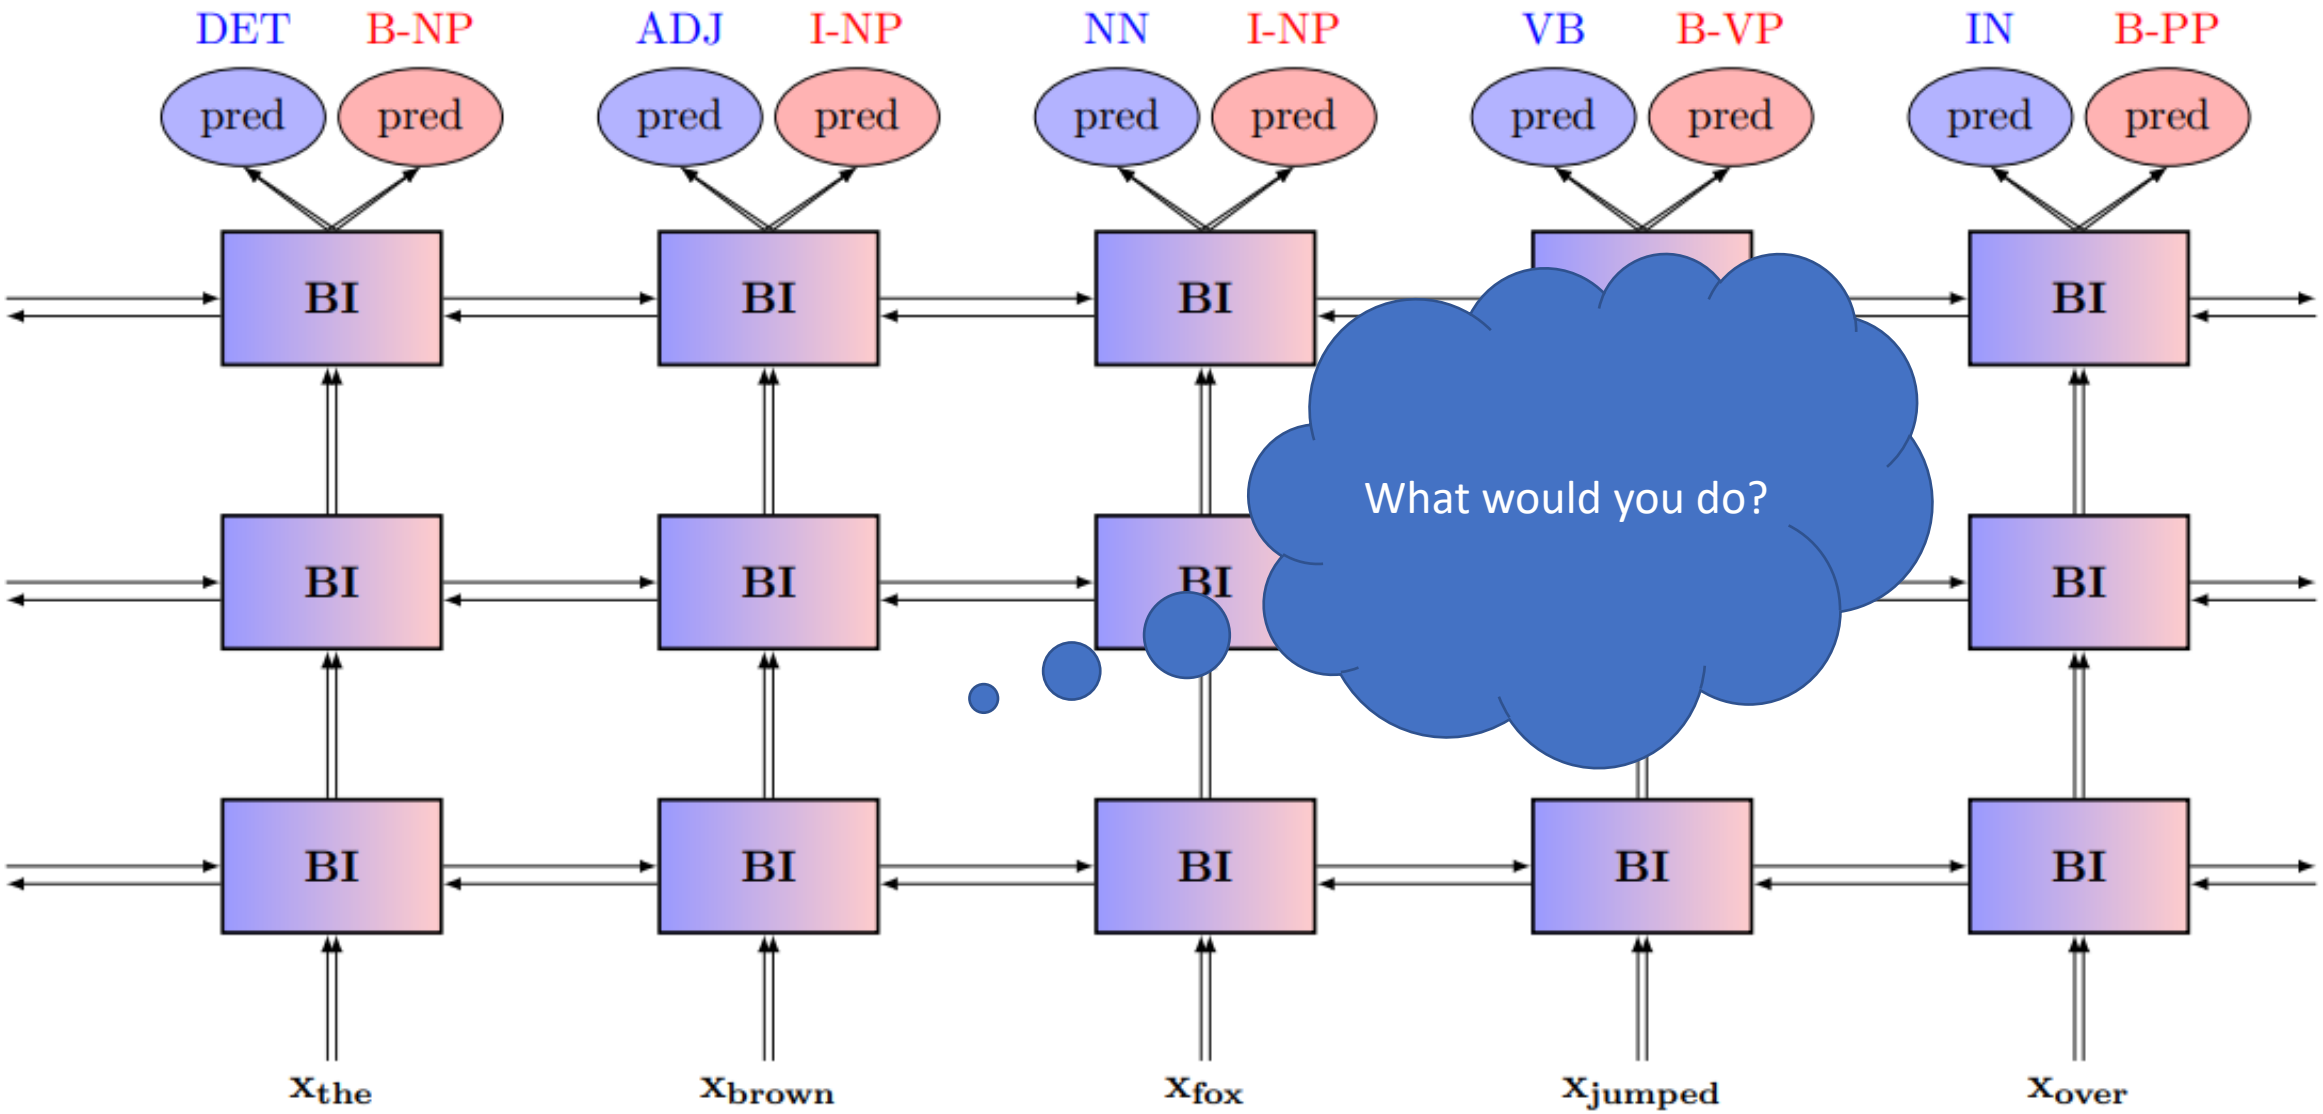
DET
B-NP
ADJ
I-NP
NN
I-NP
VB
B-VP
IN
B-PP
pred
BI
What would you do?
x_{the}
x_{brown}
x_{fox}
x_{jumped}
x_{over}

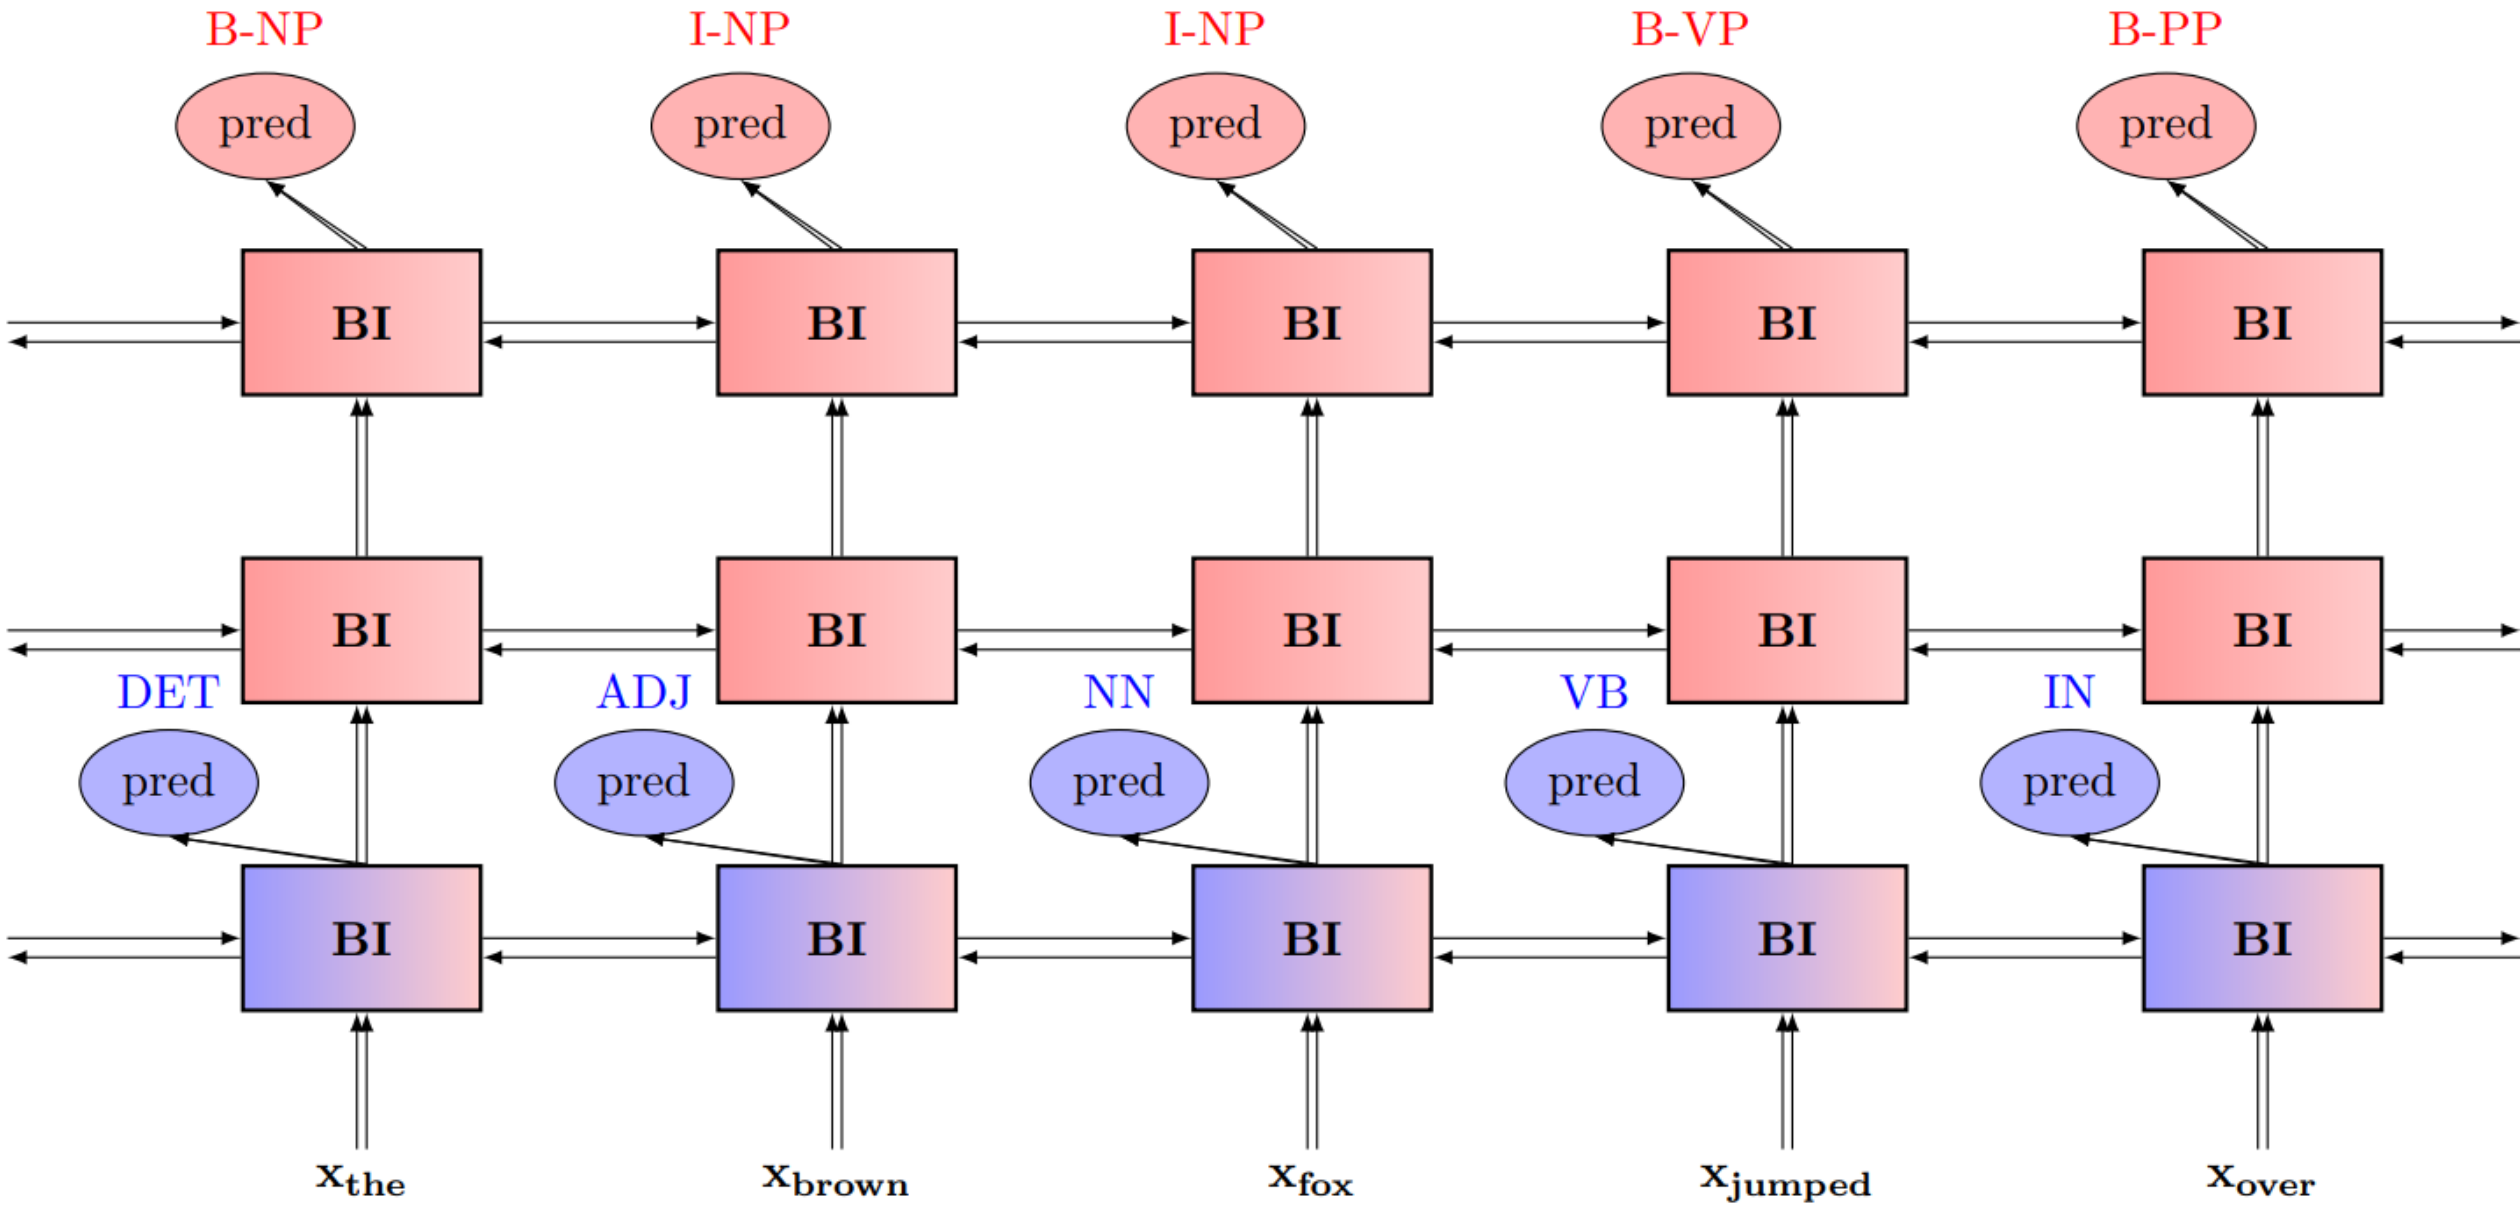
B-NP
I-NP
I-NP
B-VP
B-PP
pred
BI
DET
ADJ
NN
VB
IN
x_{the}
x_{brown}
x_{fox}
x_{jumped}
x_{over}

RNN - Summary

- Many tasks can be solved with the RNN family
  - And many ARE being solved
- Be creative with the architecture